



A Quick Primer on TensorFrames Apache Spark and TensorFlow Together

Tom Drabas
Data Scientist II

Denny Lee
Principal Program Manager

Session Goals

- What is Deep Learning?
- A primer on feature learning
- What is feature engineering?
- What is TensorFlow?
- Build a Deep Learning model to recognize objects in images
- Introducing TensorFrames
- TensorFrames – quick start

Tom Drabas

Data Scientist II, Core Data Science (WDG)

All-around geek and nerd

Close to 15 years of industry experience: analytics, operations research and data science.

Worked in airlines, telecommunications, logistics and technology

Interested in big data, algorithms, neural networks, and... vintage electronics



Denny Lee

Principal Program Manager, Azure Cosmos DB

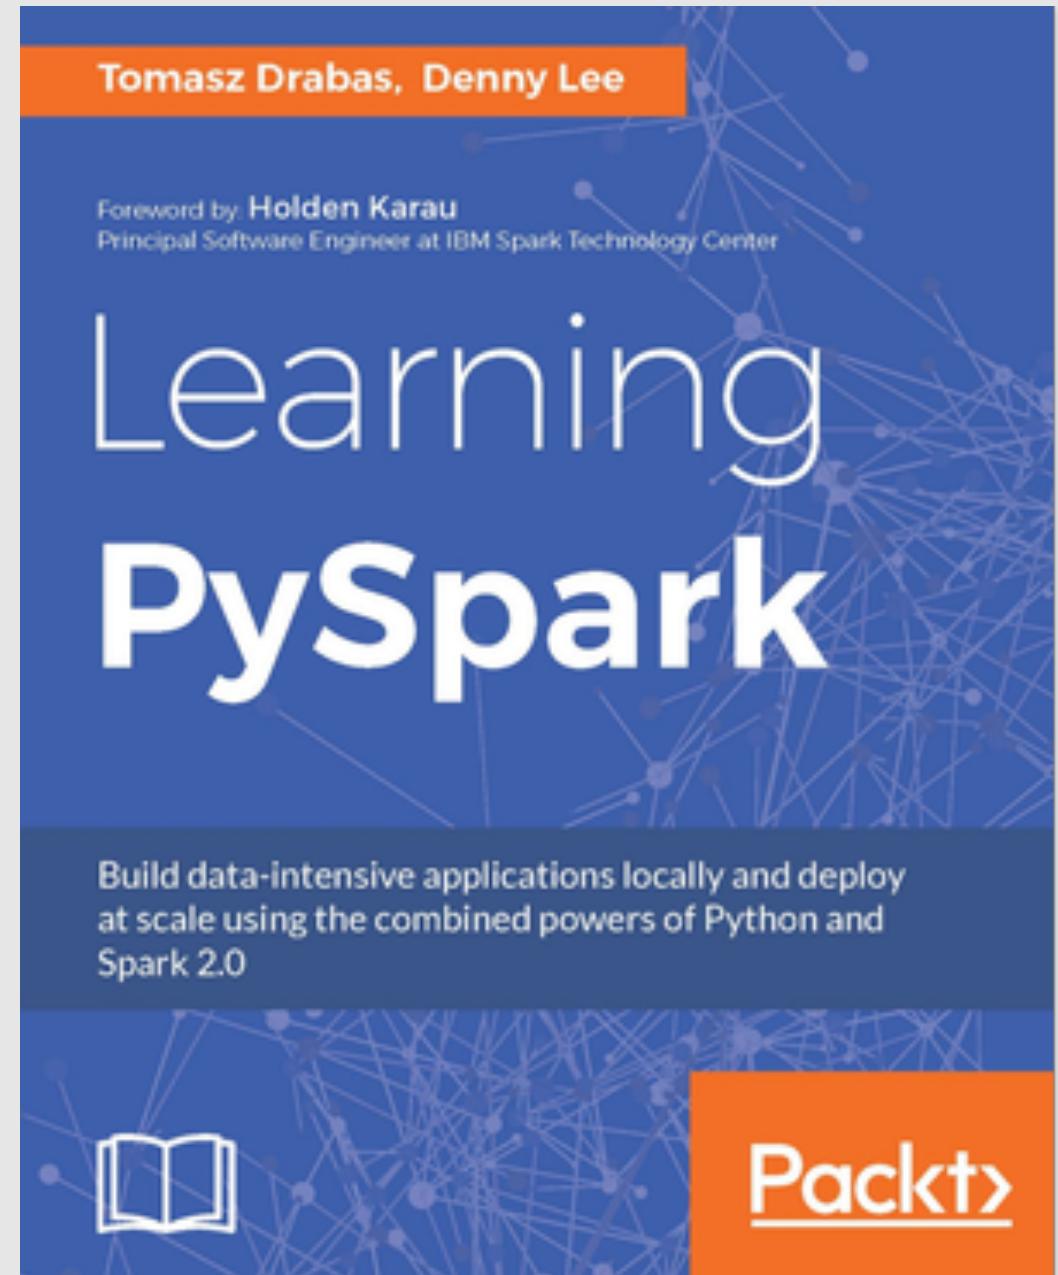
Former:

- Technology Evangelist @ Databricks
- Senior Director of Data Sciences Engineering at Concur (now part of SAP)
- Principal Program Manager at Microsoft

Hands-on Data Engineer, Architect more than 15y developer internet-scale infrastructure for both on-premises and cloud including Bing's Audience Insights, Yahoo's 24TB SSAS cube, and Isotope Incubation Team (HDInsight).



Jump Start into Python and Apache Spark



Many Thanks!



Tim Hunter
Software Engineer
Databricks

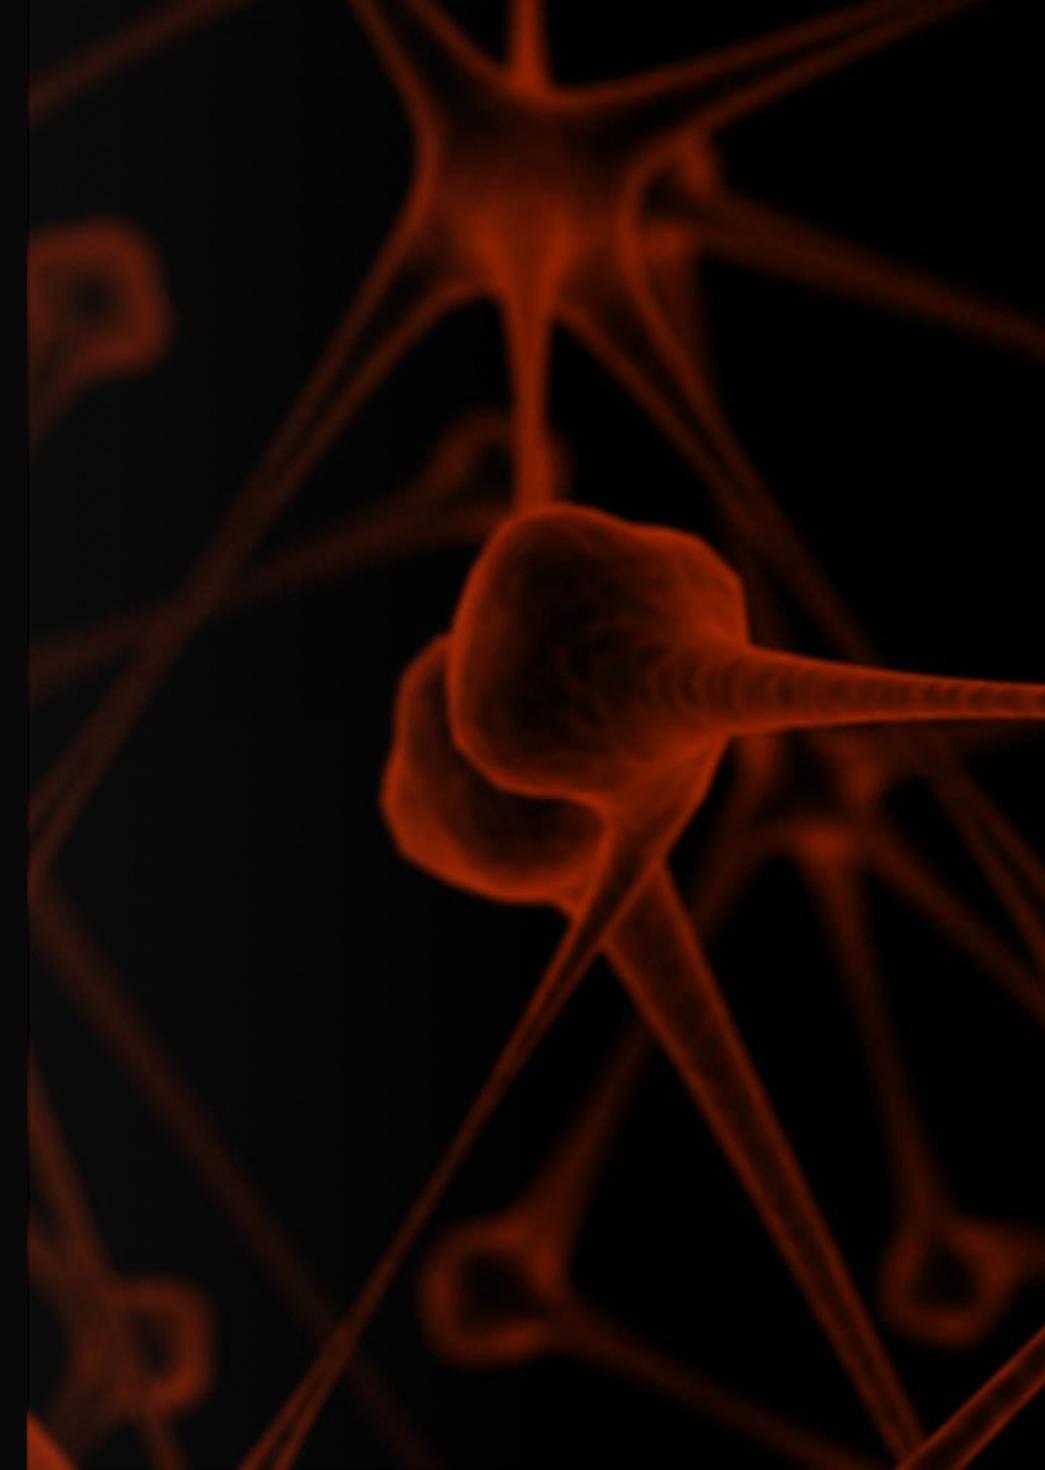


Joy Qiao,
Senior Solution Architect
Microsoft Algorithms and Data
Science Solutions



Bilal Obeidat
Solution Architect
Databricks

Prerequisites



Tutorial Prerequisites

- PySpark
 - As of Spark 2.2-ish, pip install pyspark
- Databricks Community Edition
 - <https://databricks.com/try-databricks>

Get Databricks Community Edition

<http://databricks.com/try-databricks>

Select a version to get started.

FULL-PLATFORM TRIAL

[Put Apache Spark to work](#)

- Unlimited clusters
- Notebooks, dashboards, production jobs, RESTful APIs
- Interactive guide to Spark and Databricks
- Deployed to your AWS VPC
- BI tools integration
- 14-day free trial (excludes AWS charges)

[START TODAY](#)

COMMUNITY EDITION

[Learn Apache Spark](#)

- Mini 6GB cluster
- Interactive notebooks and dashboards
- Public environment to share your work

[START TODAY](#)

Sign Up for Databricks Community Edition



Sign Up for Databricks Community Edition

First Name *

Doctor

Last Name *

Who

Company Name *

Tardis

Work Email *

put-your-email-here

Password *

Confirm Password *

Phone Number

425-555-1212

What is your intended use case? *

Personal - Learning Spark

How would you describe your role? *

Data Scientist



I'm not a robot



Sign Up

Confirm your email address



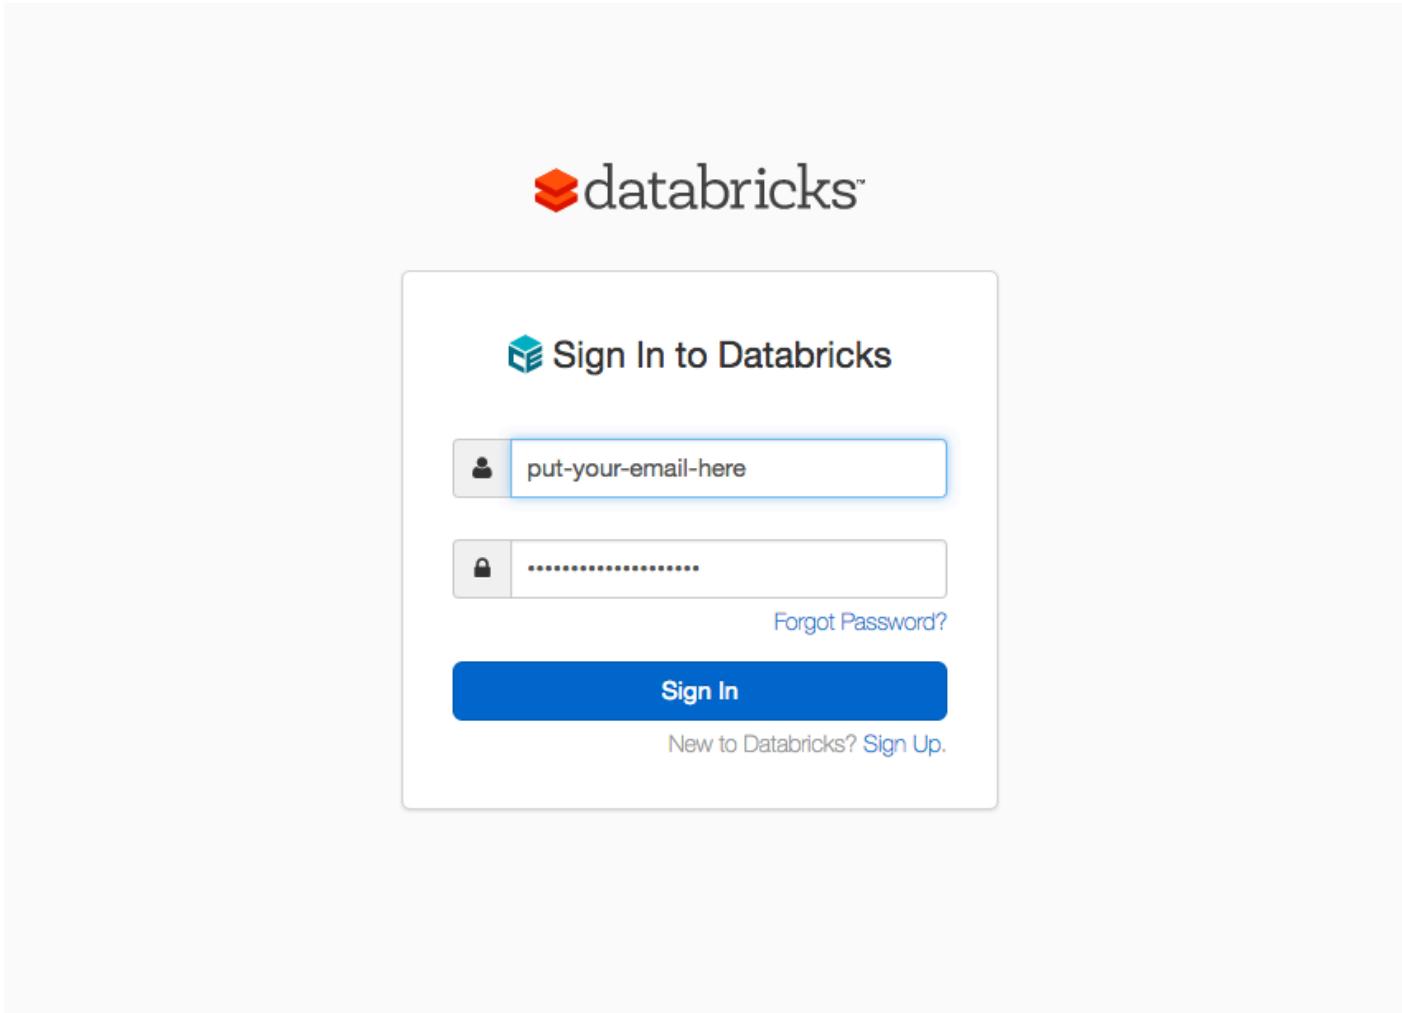
Please Confirm Your Email Address

You will receive an email with a link to confirm your email address. Please click the link to complete the signup process. **If you haven't received the email, please check your spam folder.**

Contact feedback@databricks.com if you have any questions.

Sign in!

<https://community.cloud.databricks.com>



What are artificial neural networks?

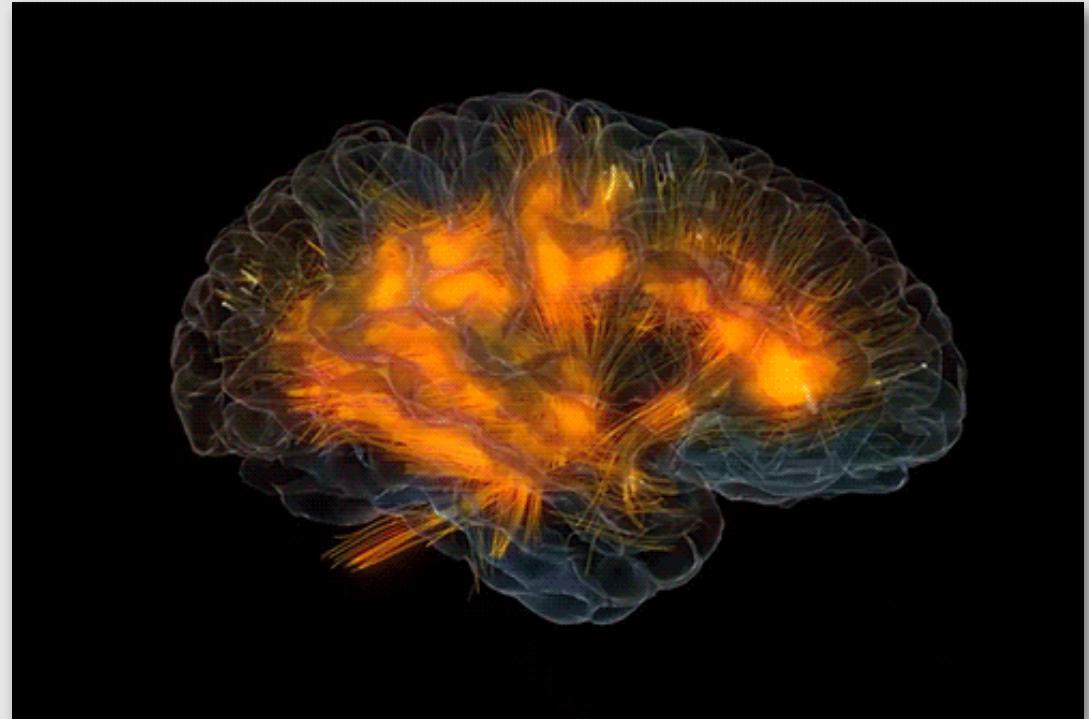


source: <http://bit.ly/2jjzFyt>

Artificial Neural Networks

An imitation of human's brain

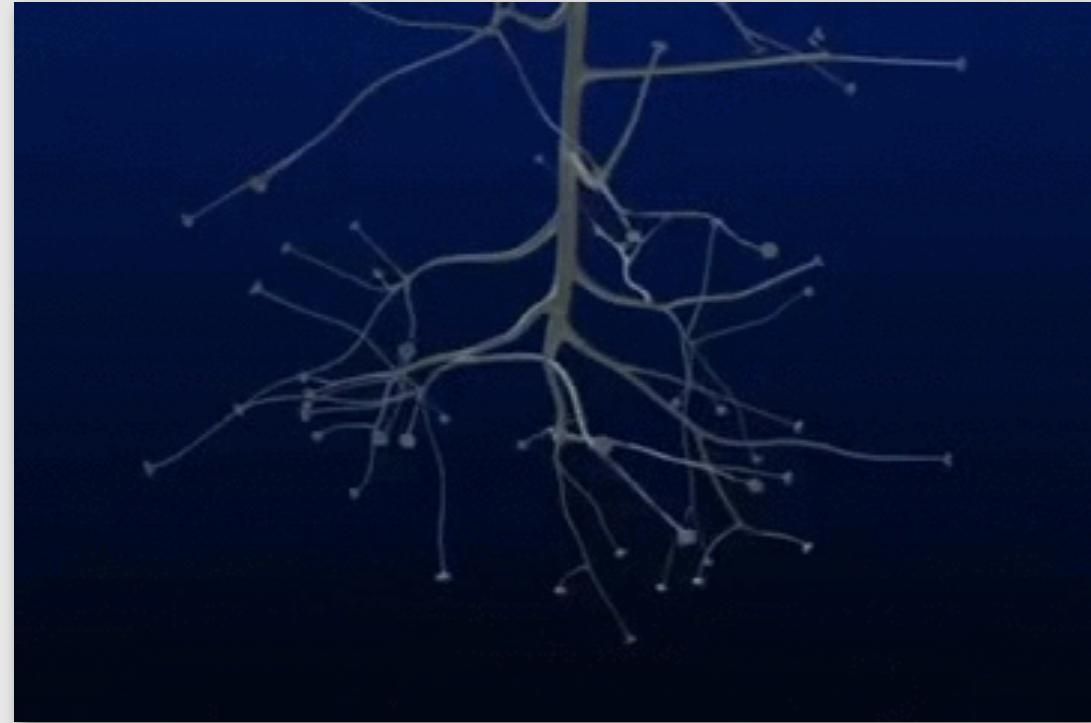
- Dense net of simple structures
- Around 100 billion neurons
- Each connected to ~10k other neurons
- 10^{15} synaptic connections



source: <http://bit.ly/2jc9mYw>

A neuron at work

- *Dendrites* receive signals
- Neuron's *cell body* acts as an accumulator
- If energy level in the neuron's body exceeds certain level it fires a short pulse through the *axon* ended with *synaptic terminals*



source: <http://bit.ly/2jjGINd>

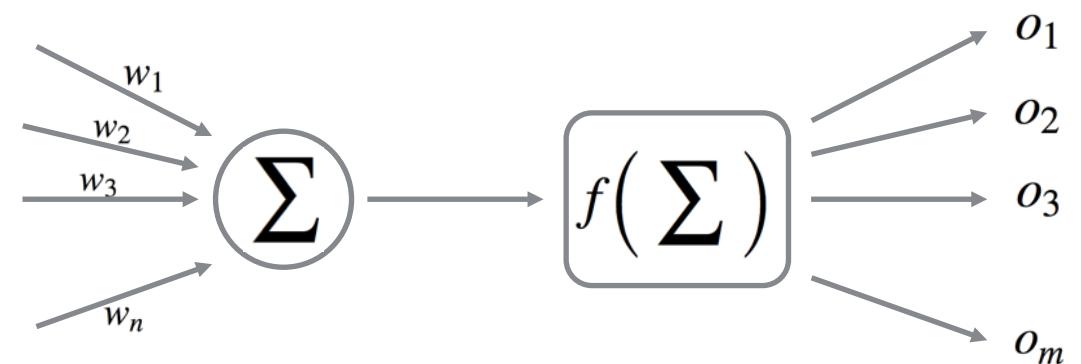
Conceptual mathematical model

- Receives input from n sources
- Computes weighted sum

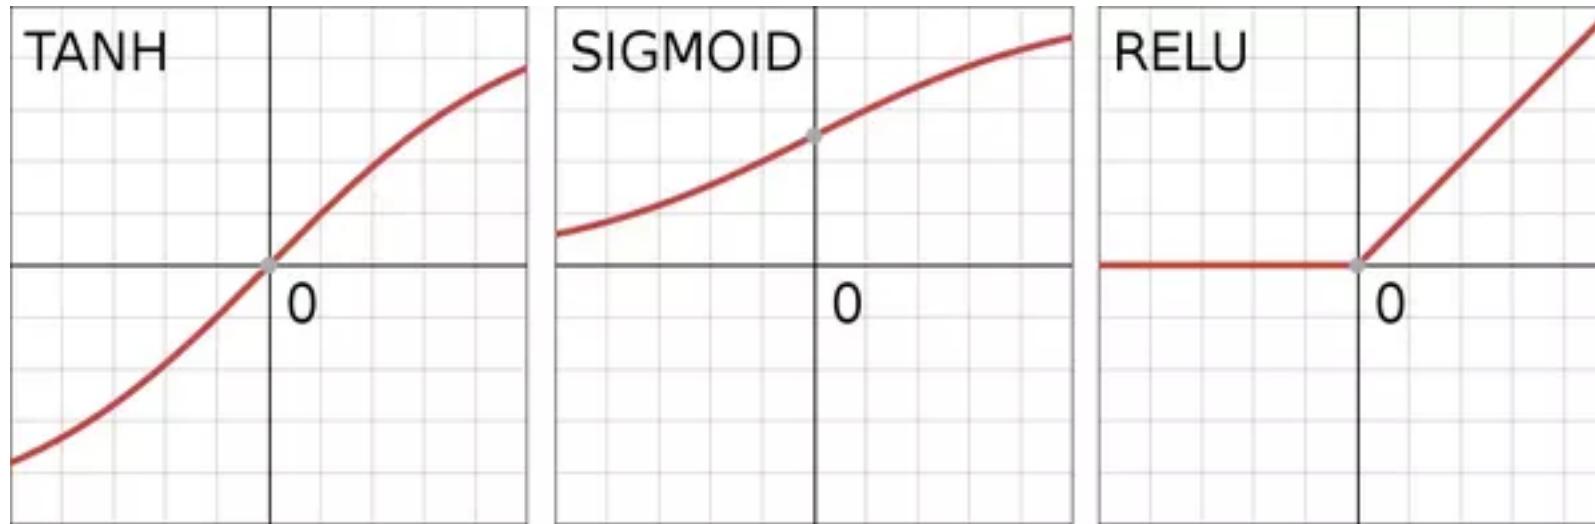
$$h_1 = x_1 w_1 + x_2 w_2 + \dots + x_n w_n$$

- Passes through an *activation function*

- Sends the signal to m succeeding neurons



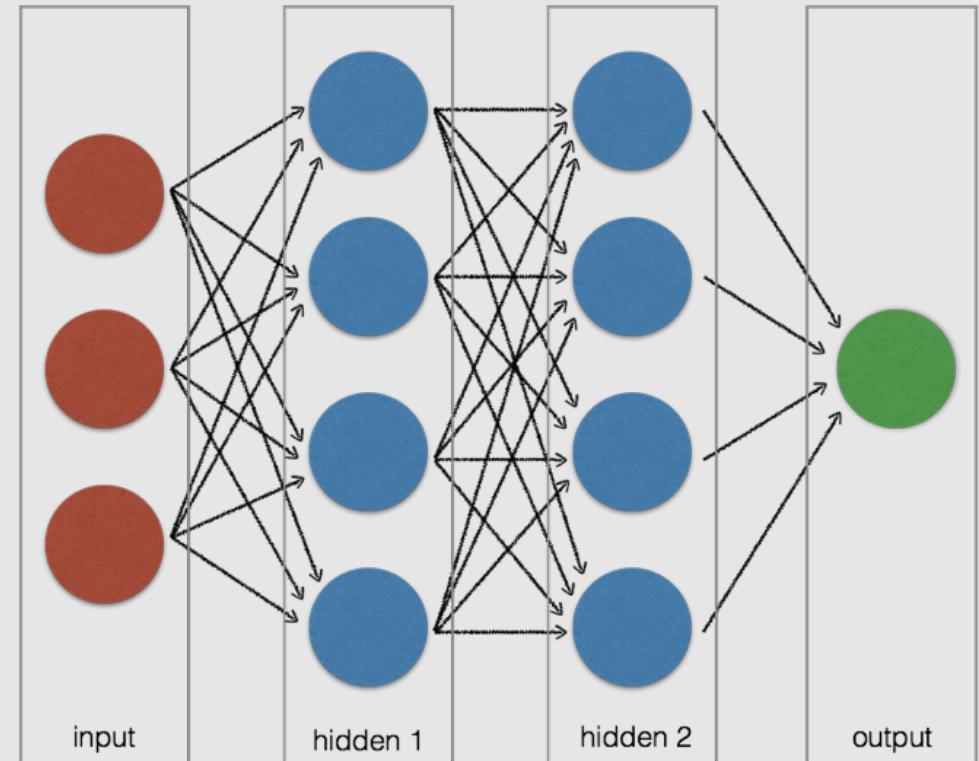
Activation functions



- Bias (threshold) activation function was proposed first
- Sigmoid and tanh introduce non-linearity with different codomains
- ReLu is the latest and greatest due to backprop (more later)

Artificial Neural Network

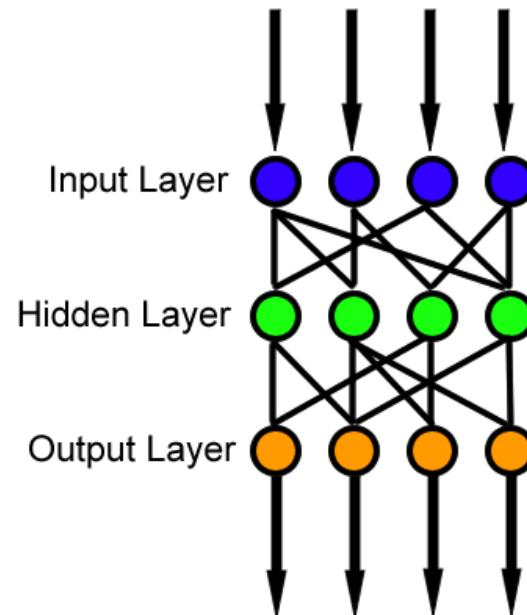
- A layer is a set of neurons
- Typically consists of input, hidden and output layers
- Input layer acts as an interface
- The neurons in hidden layer(s) and the output modify data
- A black-box model



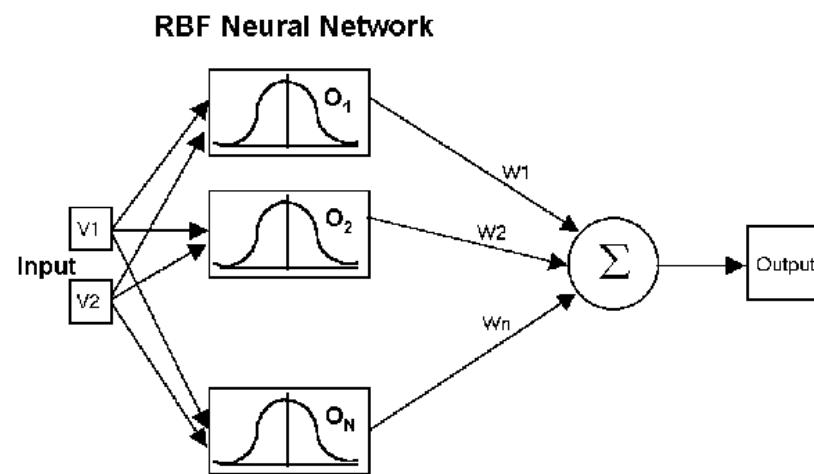
Types of neural networks

Some examples

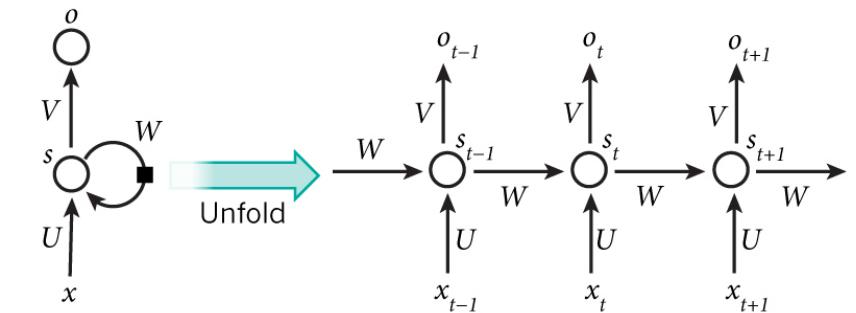
Single- or multi-layer
feedforward networks



Radial Basis Function



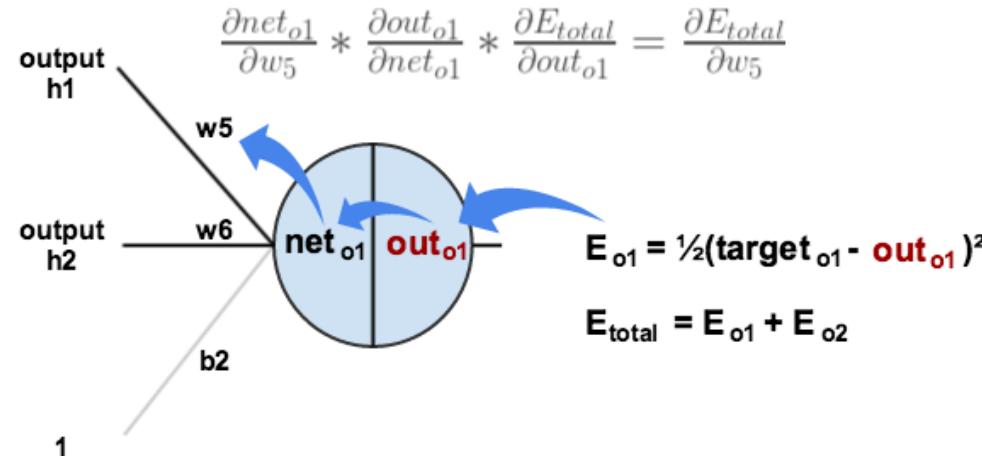
Recurrent Neural Networks



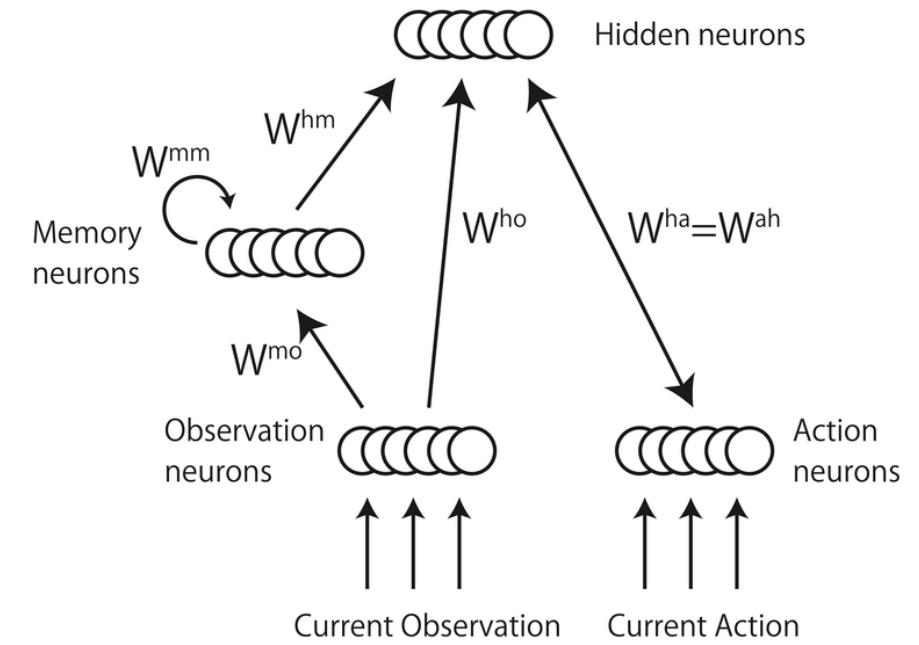
Learning

Some examples

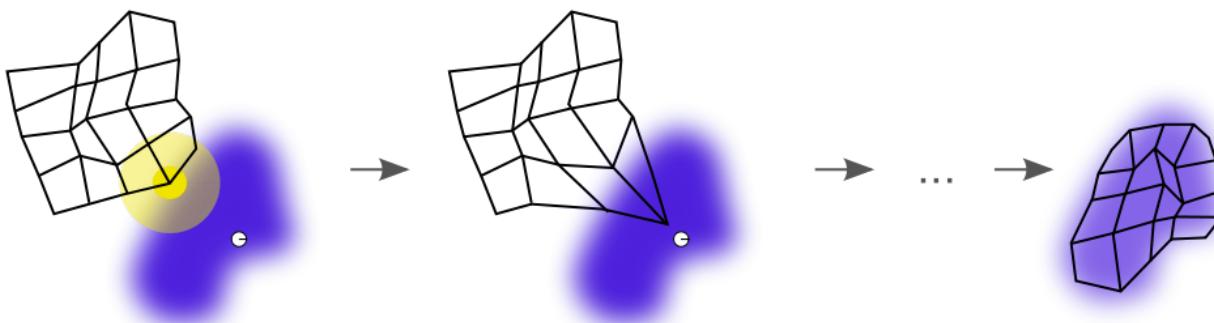
Backpropagation



Memory based learning

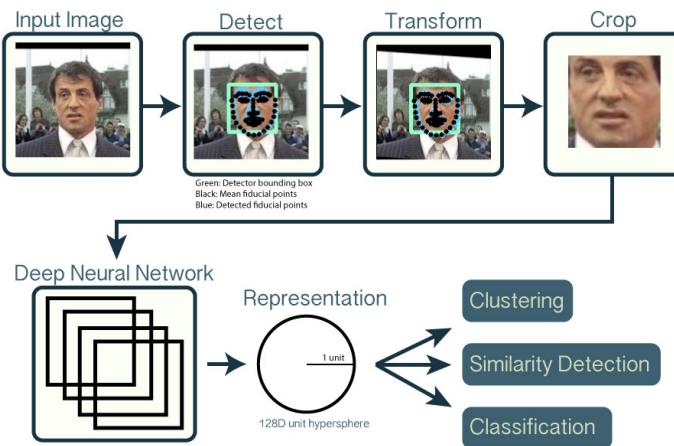


Self-organizing maps (competitive learning)

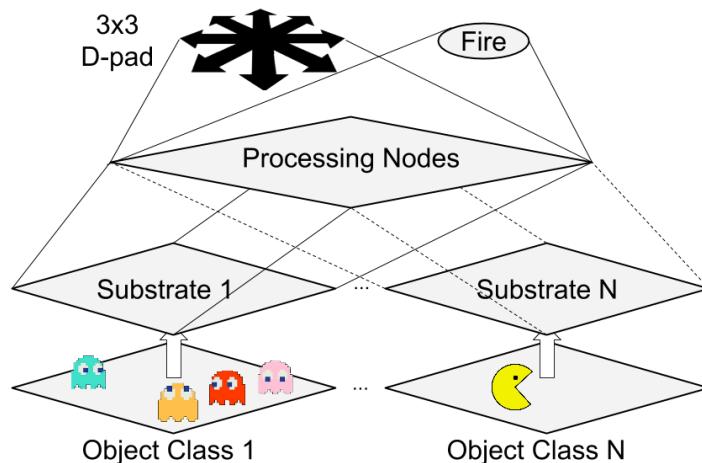


Applications of neural networks

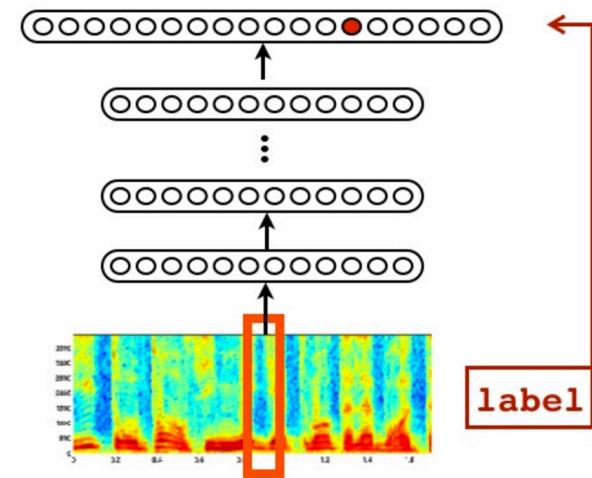
Object Recognition: Facial



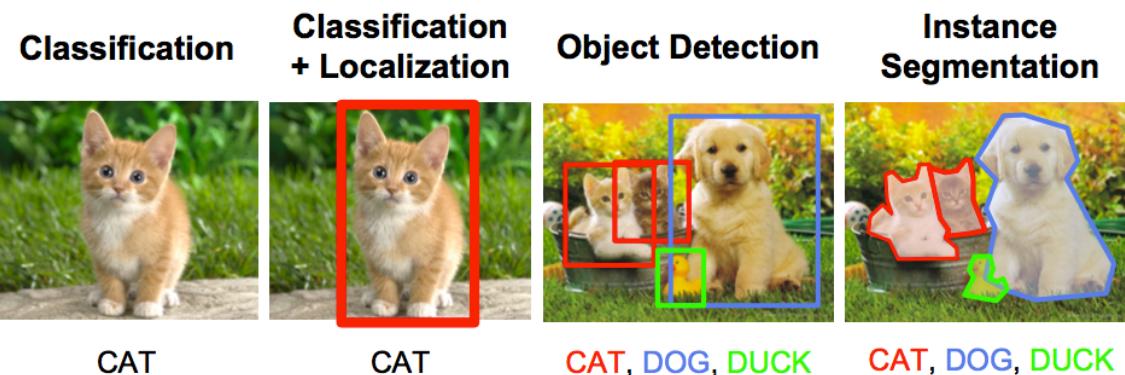
Game Playing



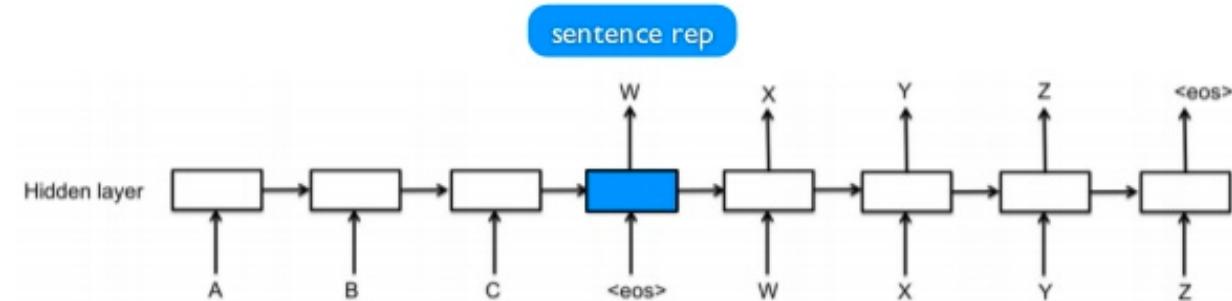
Object Recognition: Speech



Object Classification



Language Translation



What is Deep Learning?



source: <http://bit.ly/2jcquh5>

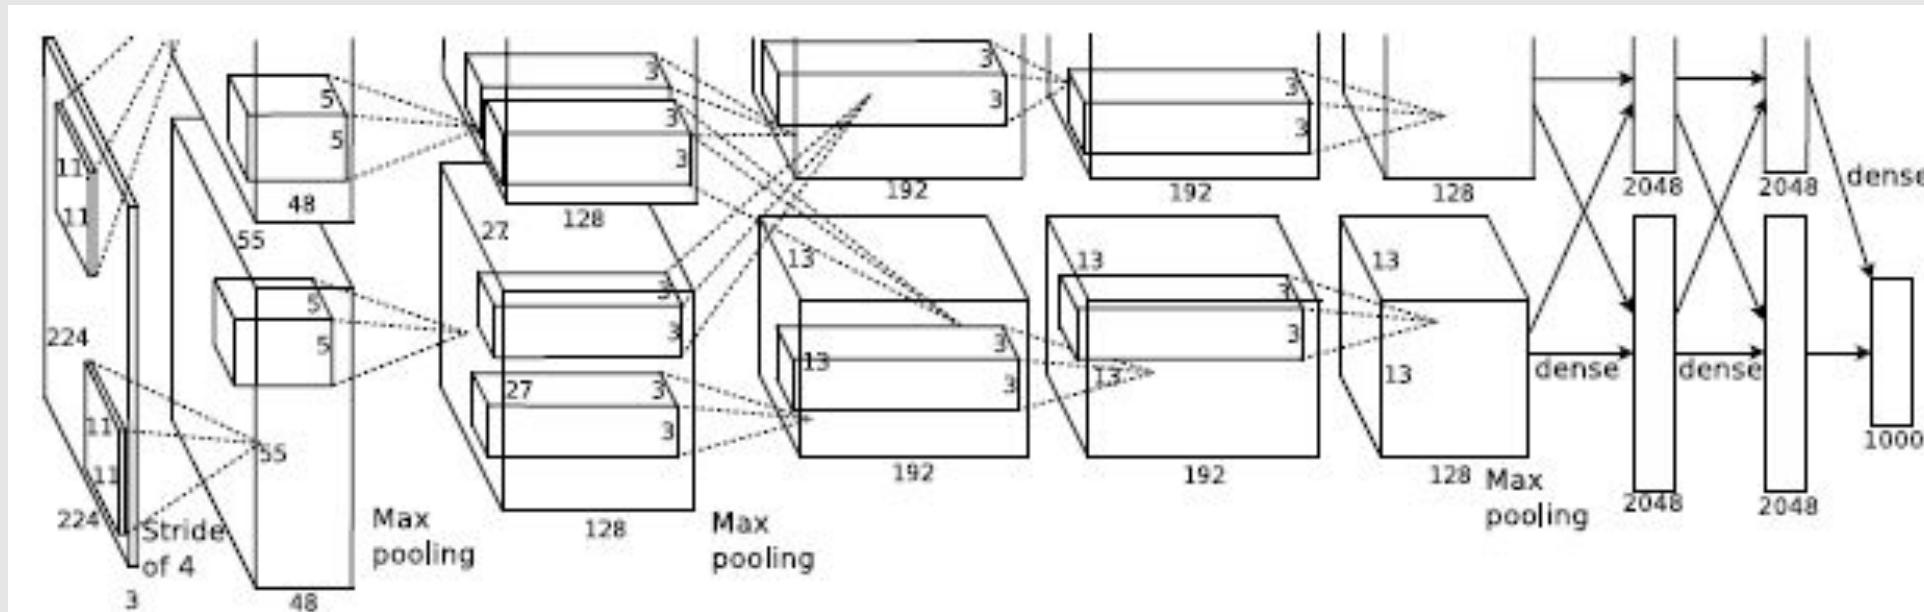
Deep Learning

The next step in AI

Deep learning

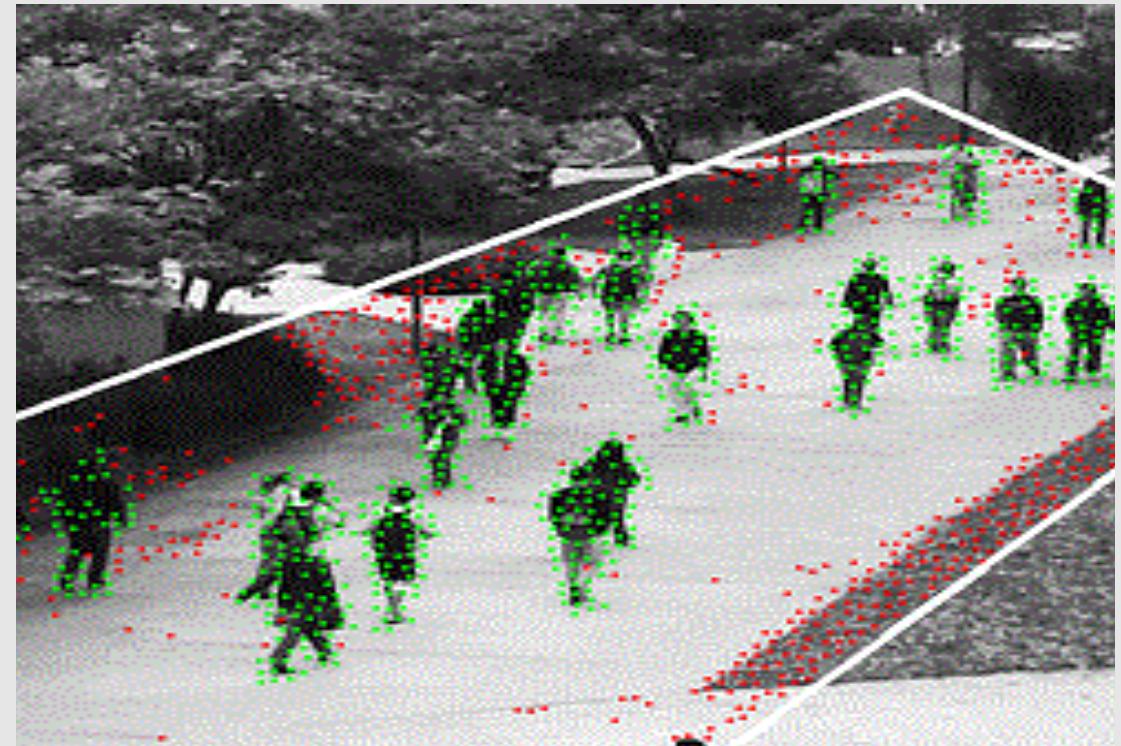
- Vast and sophisticated ANN structure
- Tens, hundreds, thousands... of hidden layers
- Ability to not only map the input to the output but recognize what inputs are needed for the task

Example: AlexNet (ImageNet 2012 winner); Source: <http://bit.ly/2tNa3Ps>



(Automated) Feature learning

- Determine which features to use for a given task
- Ability to (implicitly) perform feature transformations and dimensionality reduction
- Automated feature engineering and learning eliminates the need for human intervention

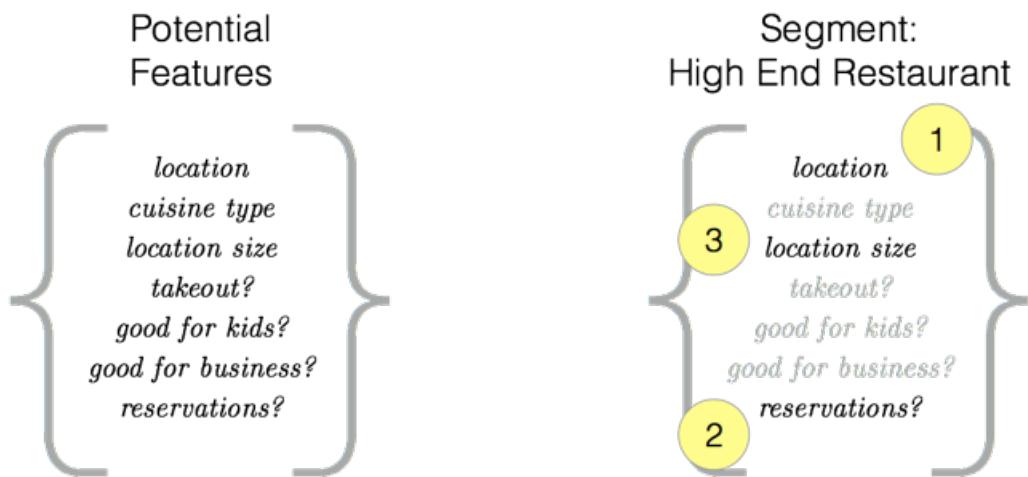


source: <http://bit.ly/2jwytoN>

In practice

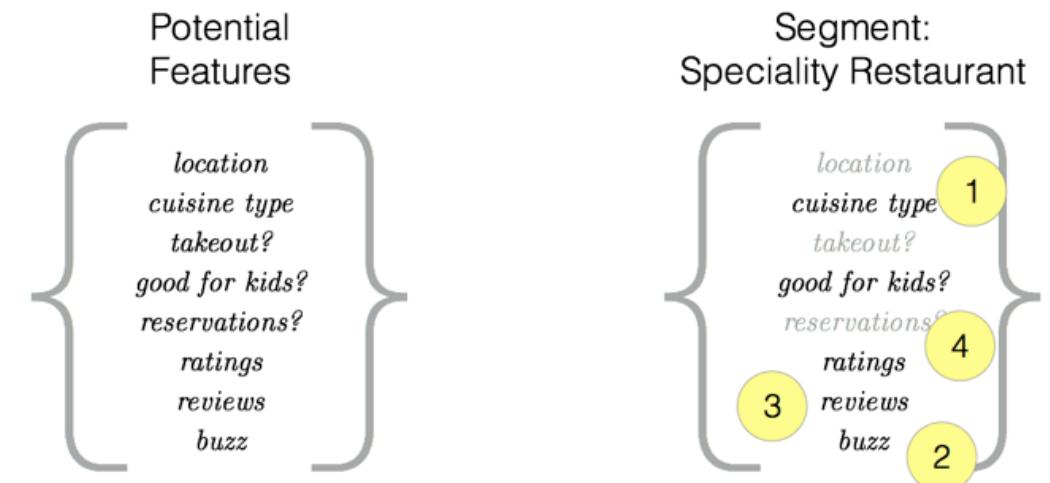
High-end restaurant

Key features are often related to location, ability to make reservations for large parties, and diversity of wine list

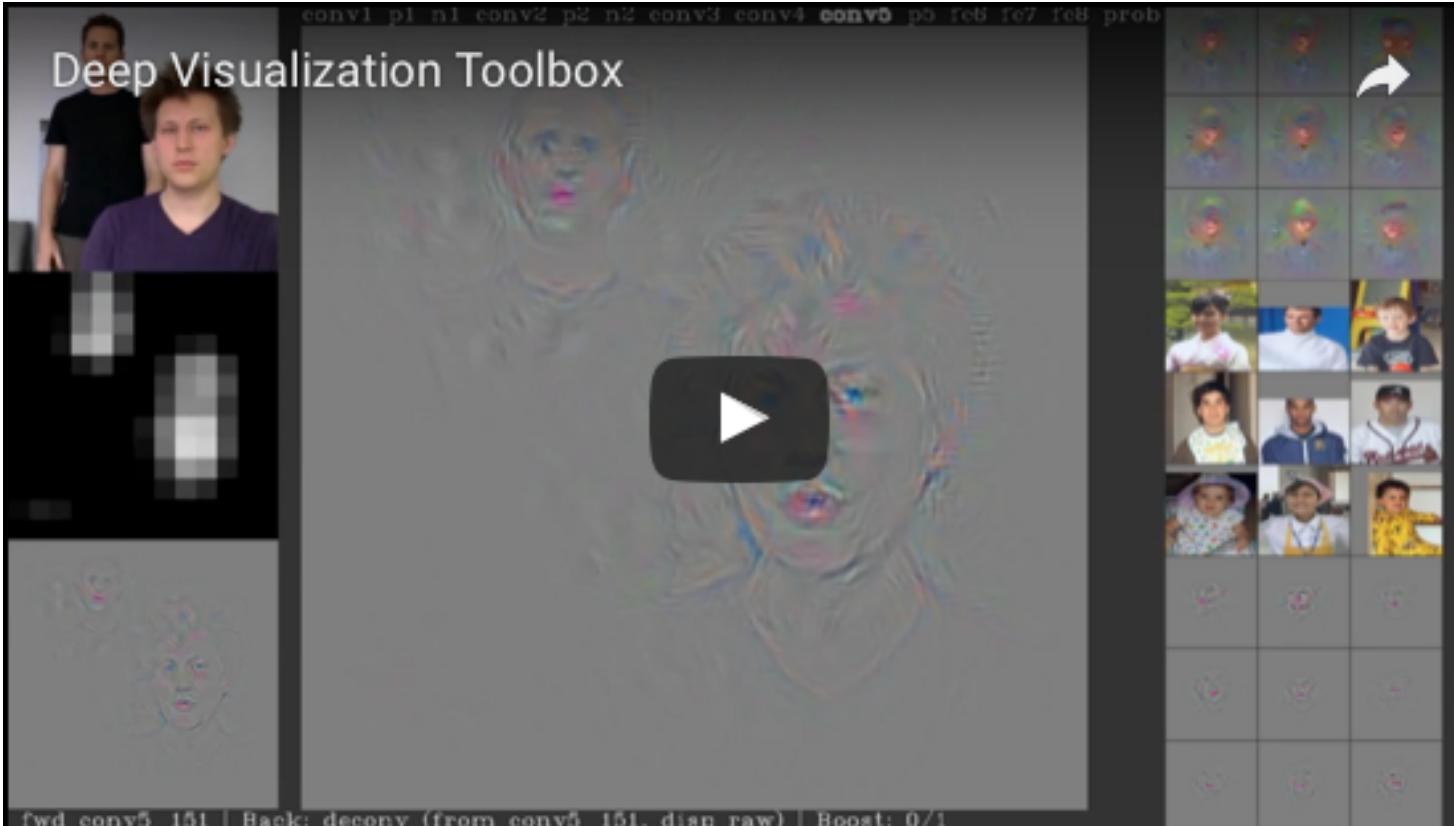


Specialty restaurant

Focus is on the reviews, ratings, social media buzz, and possibly if the restaurant is good for kids



Deep Visualization Toolbox



Demonstrate different layers of the conv net neuron activation

<https://github.com/yosinski/deep-visualization-toolbox>

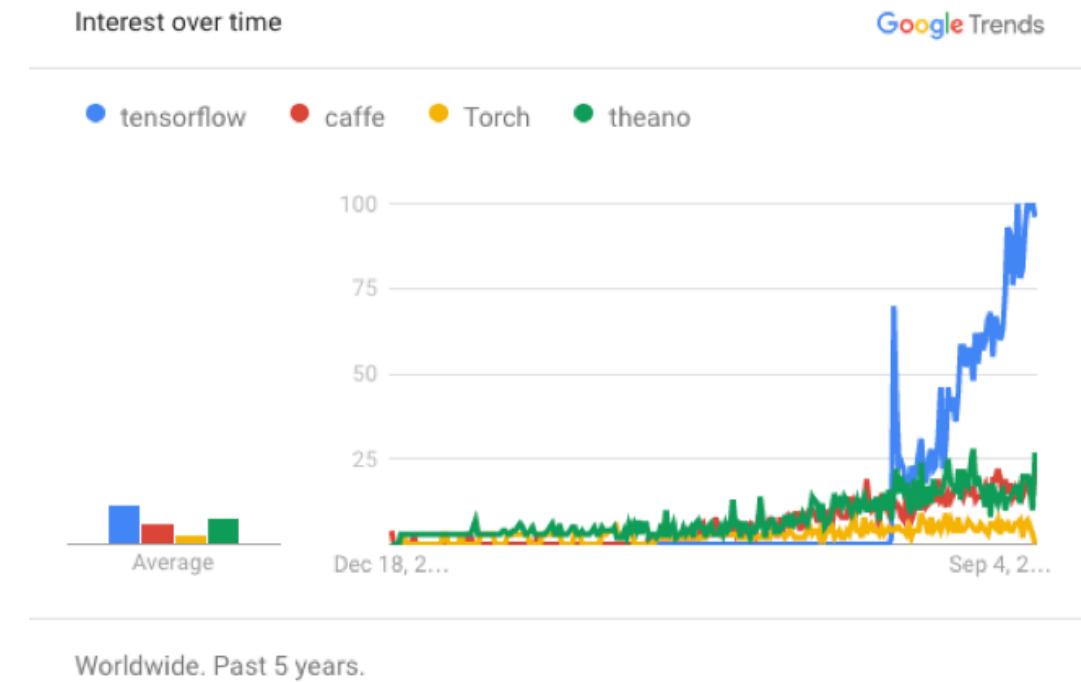


TensorFlow

Technology from Google

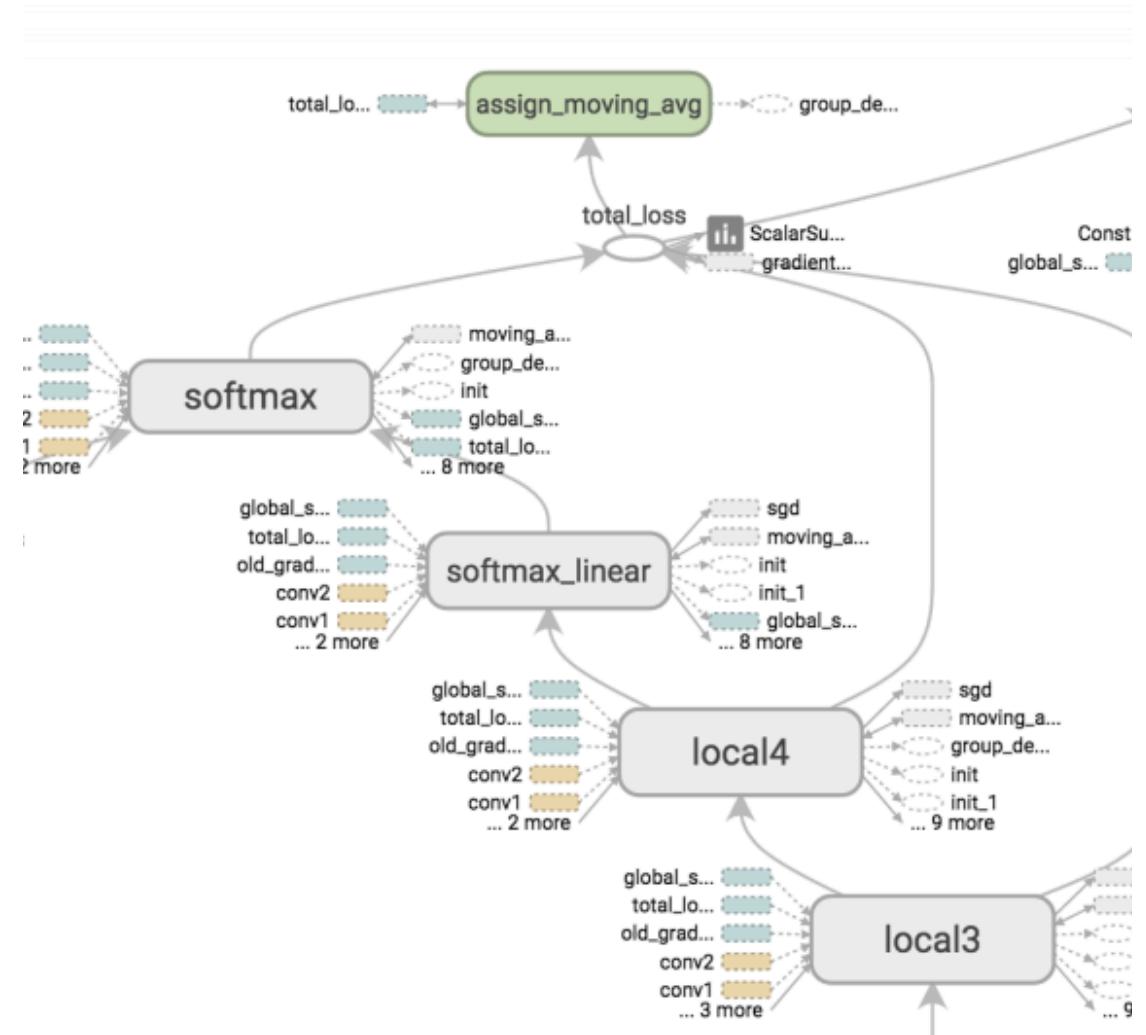
What is TensorFlow

- Open source software library for numerical computation using data flow graphs
 - Based loosely on neural networks
 - Built in C++ with a Python interface
 - Quickly gained high popularity
 - Supports GPU computations



Data Flow Graph

- Nodes represent mathematical operations
 - Edges are multidimensional arrays (tensors)
 - Thus, TensorFlow is a flow of multidimensional arrays (tensors) between the mathematical operations (nodes)
 - i.e. flow of *tensors*



Concept

- t_1 is a 2×3 matrix
 - t_2 is a 3×2 matrix
 - The op_1 is a multiplication operator
-
- The above can be represented in a matrix form

$$t_1 = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \quad op_1 \quad t_2 = \begin{bmatrix} u & v \\ w & x \\ y & z \end{bmatrix}$$

$$op_1 = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} * \begin{bmatrix} u & v \\ w & x \\ y & z \end{bmatrix}$$

Example (pure Python)

1. Import TensorFlow

```
In [1]: import tensorflow as tf
```

2. Prepare the tensors and operator

```
In [2]: t1 = tf.placeholder(tf.float32)
t2 = tf.placeholder(tf.float32)

op = tf.matmul(t1, t2)
```

3. Feed the data

```
In [3]: m1 = [[3., 2., 1.]]
m2 = [[-1.], [2.], [1.]]

with tf.Session() as s:
    print(s.run([op], feed_dict={t1: m1, t2: m2}))
```

4. And get...

```
[array([[ 2.]], dtype=float32)]
```



source: <http://bit.ly/2jRN1mj>

TensorFrames (experimental)

Built for speed!

Apache Spark Engine



Scale out, fault tolerant

Python, Java, Scala, and R APIs

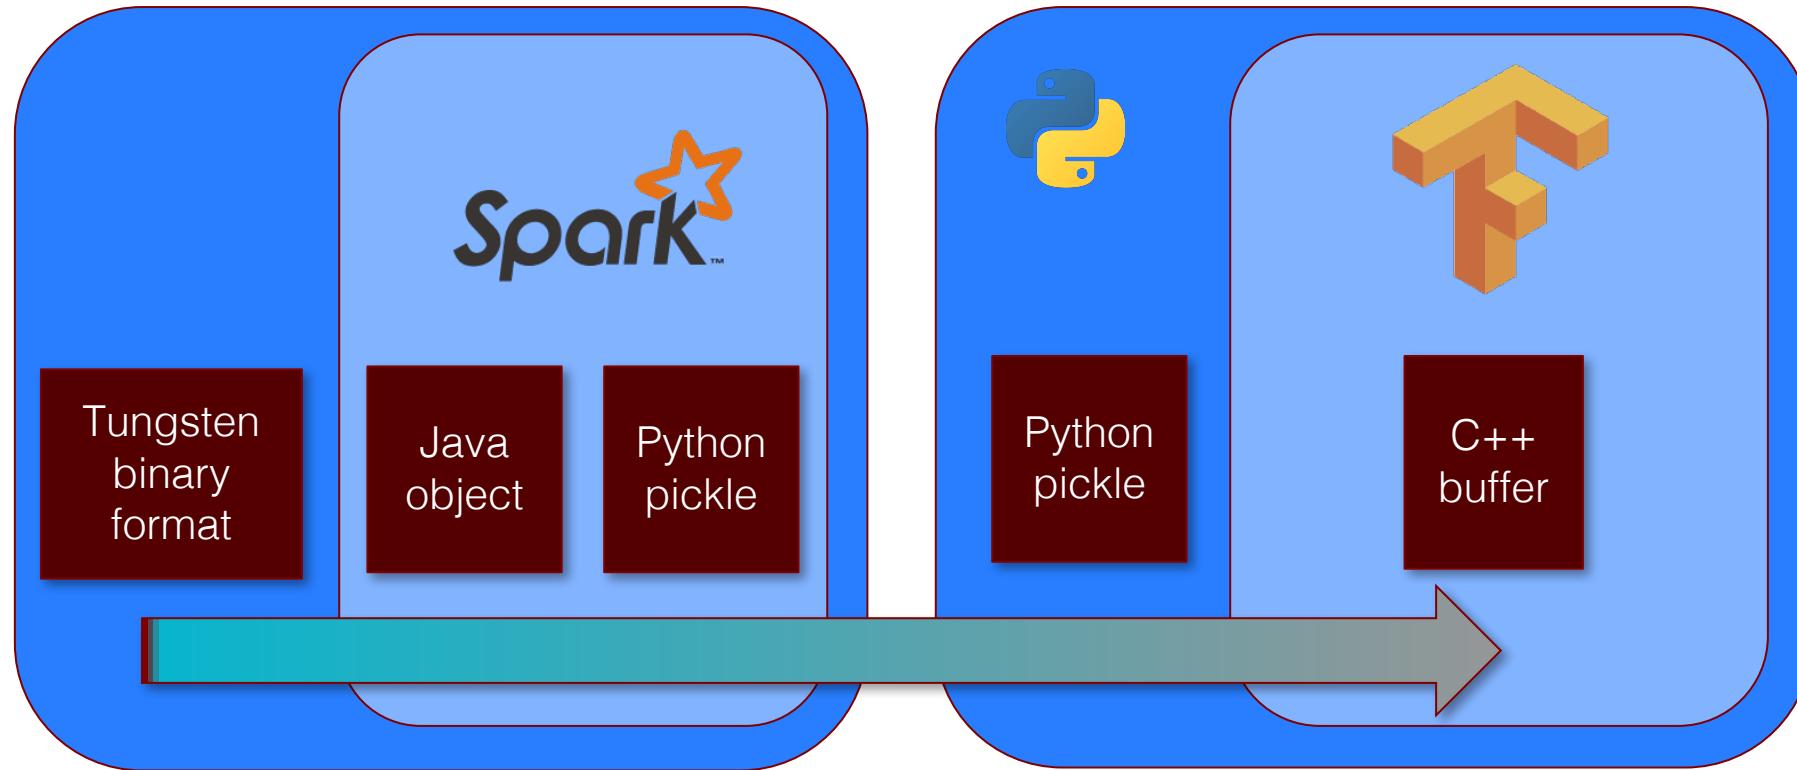
Standard libraries



...

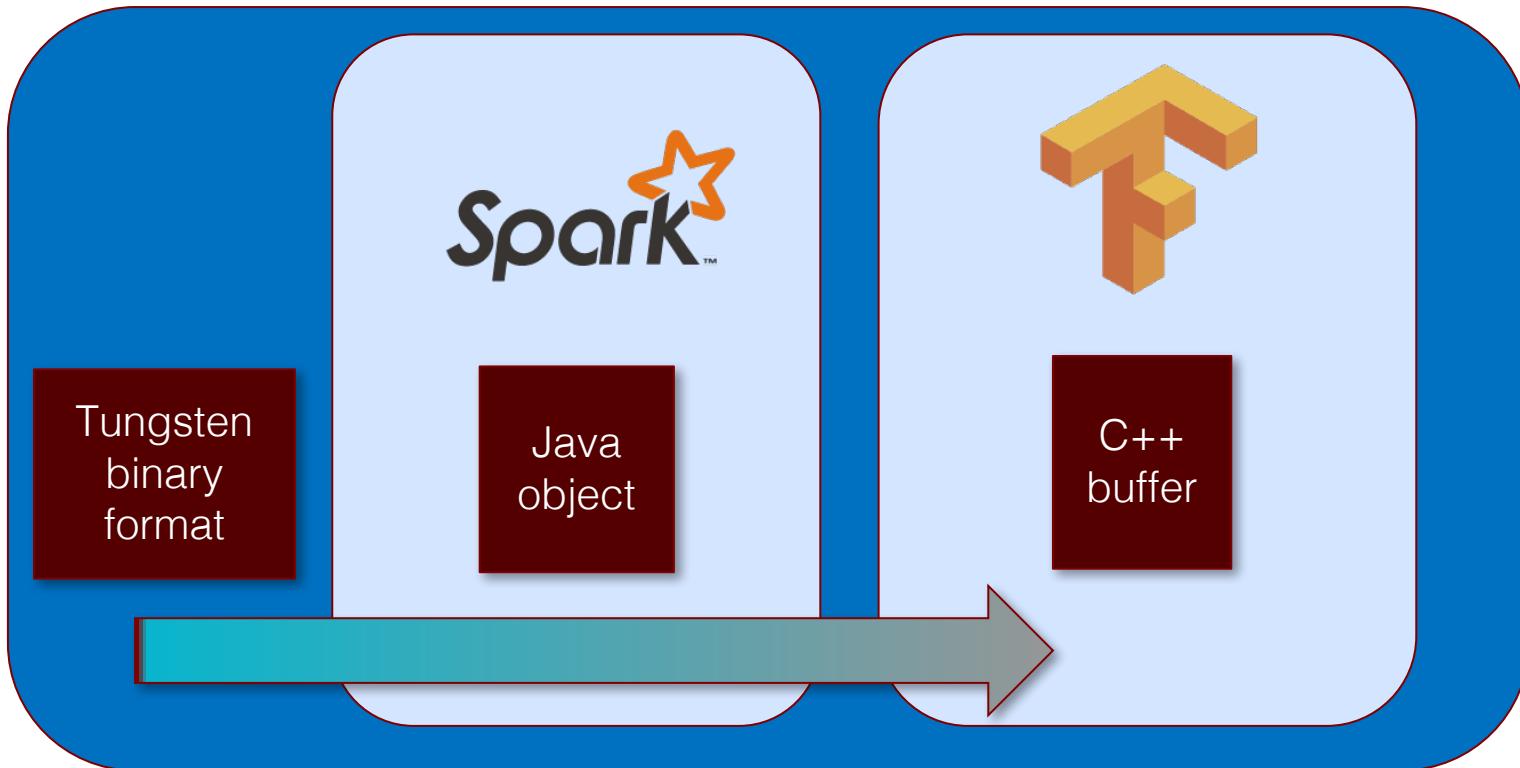
Unified engine across diverse workloads & environments

Spark to TensorFlow Communication Issues



Source: TensorFrames: Google TensorFlow with Apache Spark by Tim Hunter

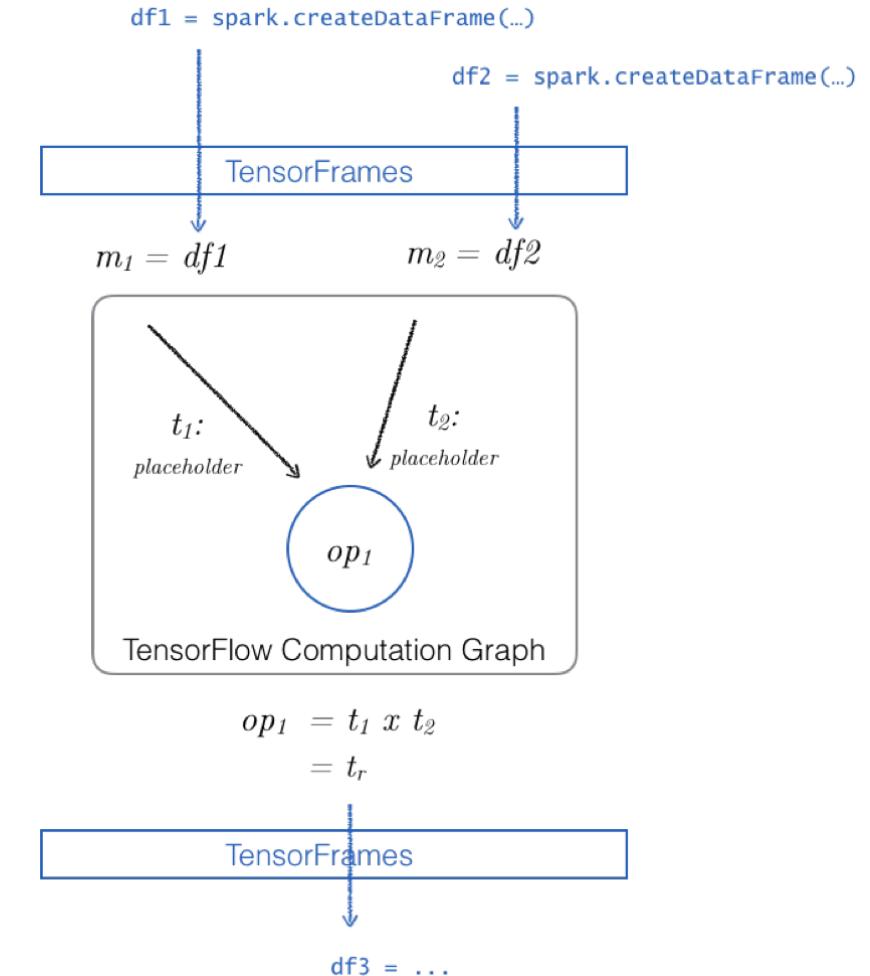
TensorFrames: Native Embedding of TensorFlow



Source: TensorFrames: Google TensorFlow with Apache Spark by Tim Hunter

An intermediary

- TensorFrames are built for speed*
- Spark DataFrames as input
- TensorFrames act as a bridge
- Can return a DataFrame

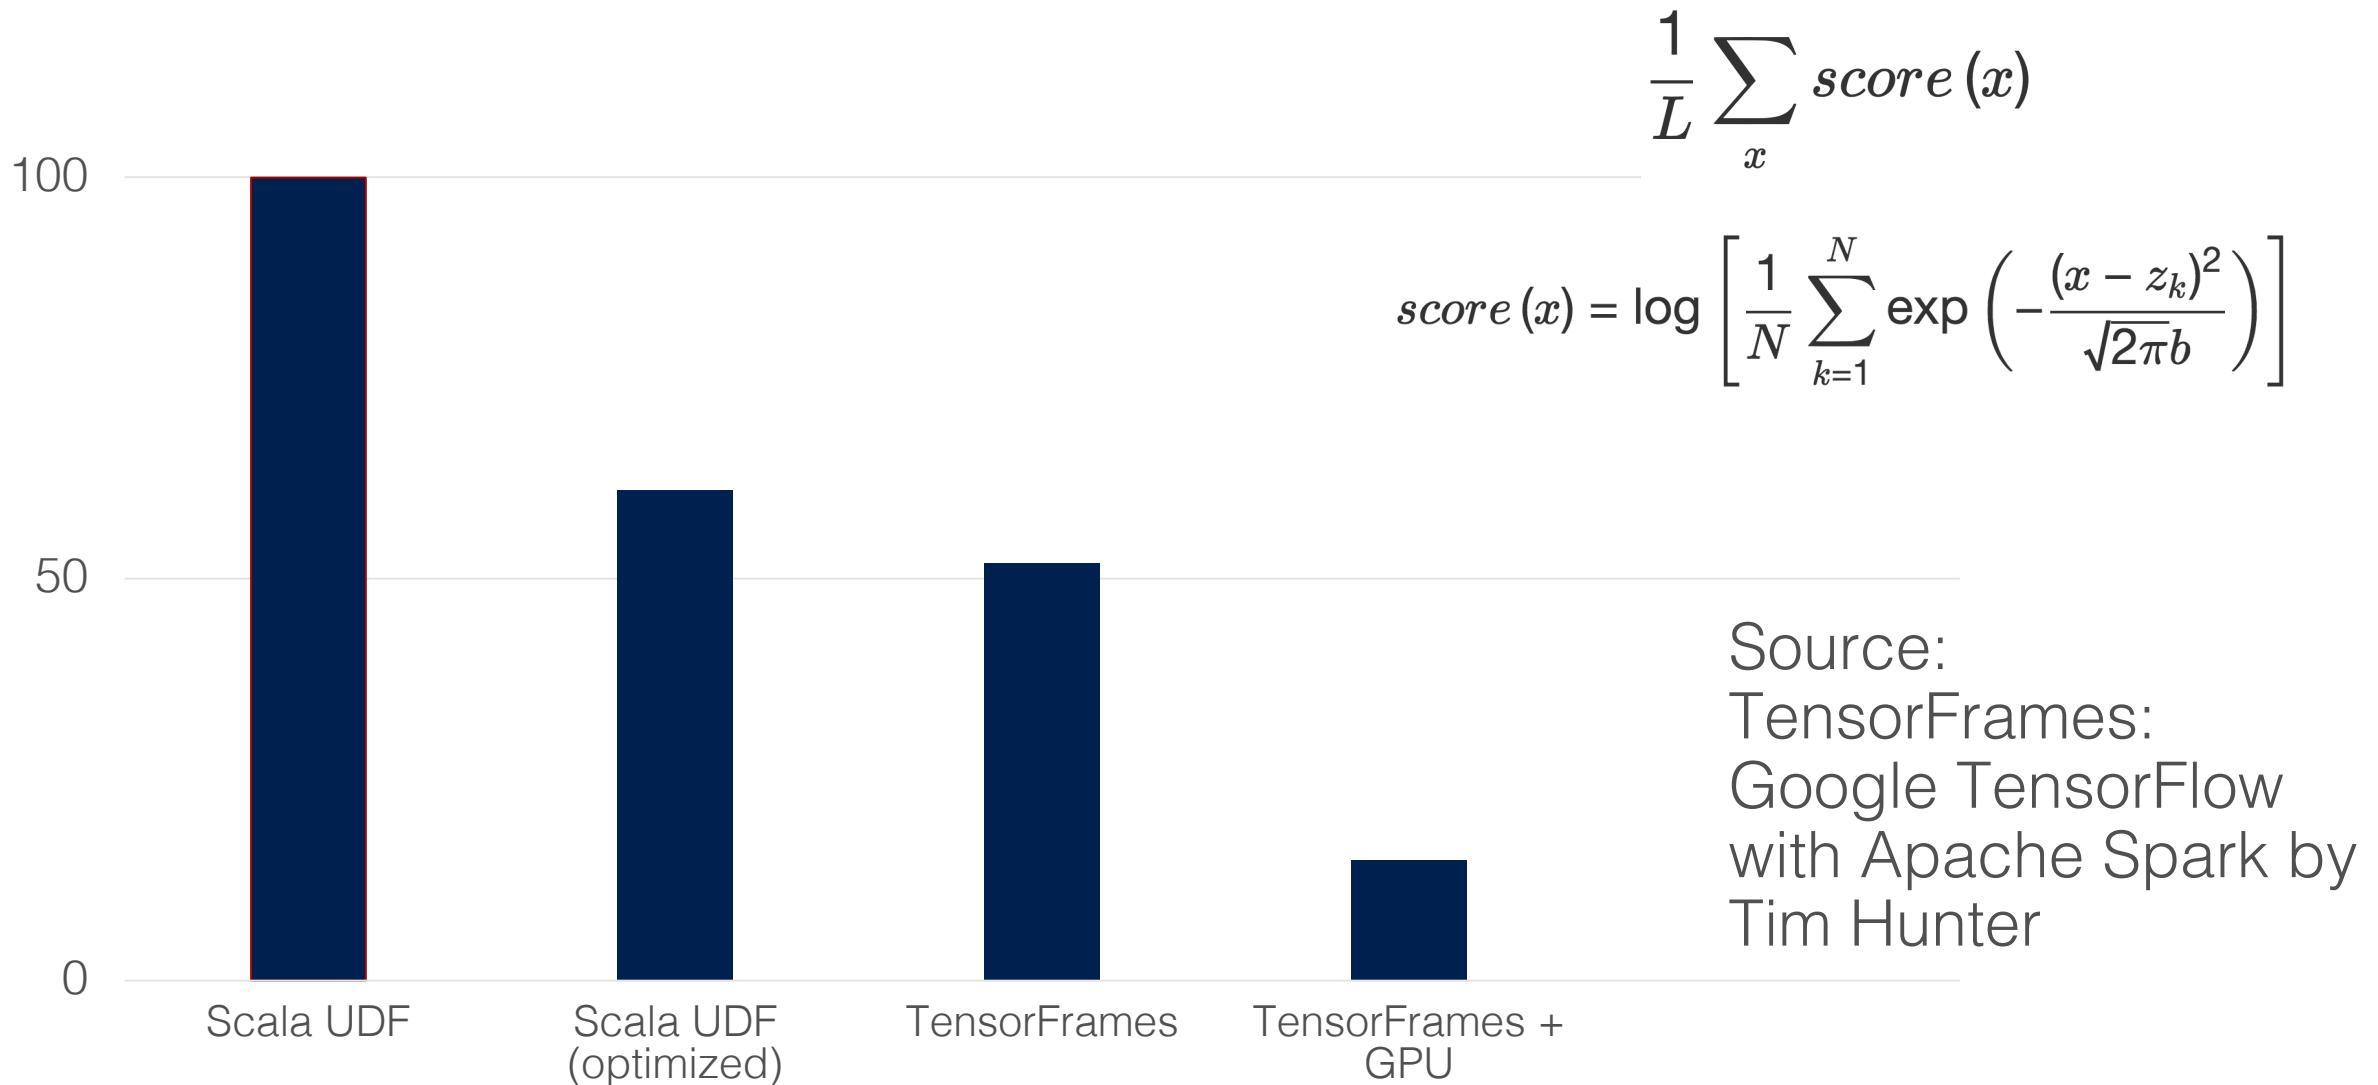


<http://bit.ly/2iVy8vm>

Why TensorFrames?

- Utilize TensorFlow with your data:
 - Allows data scientists to expand their analytics, streaming, graph, and ML capabilities to include Deep Learning via TensorFlow
- Parallel training to determine optimal hyperparameters:
 - Determine optimal learning rate: If rate is high, it will learn quickly but may not take into account of high variability
 - Number of neurons in each layer: too many results in noisy estimates, too few and the network will not learn well.

Comparing Kernel Density Scoring



Getting ready

From CLI run

```
$SPARK_HOME/bin/pyspark \  
--packages databricks:tensorframes:0.2.8-s_2.11
```

or (preferred) in Free Community Edition on Databricks (instructions shortly)

Start!

1. Import modules

2. Create some fake data

3. Create the graph

4. Block the variable

5. Add 3 to x and put in z

6. Return as df2

```
> # Import TensorFlow, TensorFrames, and Row
import tensorflow as tf
import tensorframes as tfs
from pyspark.sql import Row

# Create RDD of floats and convert into DataFrame `df`
rdd = [Row(x=float(x)) for x in range(10)]
df = sqlContext.createDataFrame(rdd)

> # Run TensorFlow program executes:
#   The 'op' performs the addition (i.e. 'x' + '3')
#   Place the data back into a DataFrame
with tf.Graph().as_default() as g:
    # The TensorFlow placeholder that corresponds to column 'x'.
    # The shape of the placeholder is automatically inferred from the DataFrame.
    x = tfs.block(df, "x")

    # The output that adds y to x
    z = tf.add(x, 3, name='z')

    # The resulting dataframe
    df2 = tfs.map_blocks(z, df)

# Note that 'z' is the tensor output from the 'tf.add' operation
print z
```

Element-wise sum and min

1. Create a DataFrame of lists
2. Analyze
3. Copy y
4. Create graph
5. Block the variables
6. Create the nodes
7. Get the results back

```
> # Build a DataFrame of vectors
  data = [Row(y=[float(y), float(-y)]) for y in range(10)]
  df = sqlContext.createDataFrame(data)

+-----+
|      y|
+-----+
| [0.0, 0.0]|
| [1.0, -1.0]|
| [2.0, -2.0]|
| [3.0, -3.0]|
| [4.0, -4.0]|
| [5.0, -5.0]|
| [6.0, -6.0]|
| [7.0, -7.0]|
| [8.0, -8.0]|
| [9.0, -9.0]|
+-----+  
  
> # Because the dataframe contains vectors, we need to analyze it first to find the
   # dimensions of the vectors.
  df2 = tfs.analyze(df)

> # Note: First, let's make a copy of the 'y' column. This will be very cheap in Spark 2.0+
  df3 = df2.select(df2.y, df2.y.alias("z"))

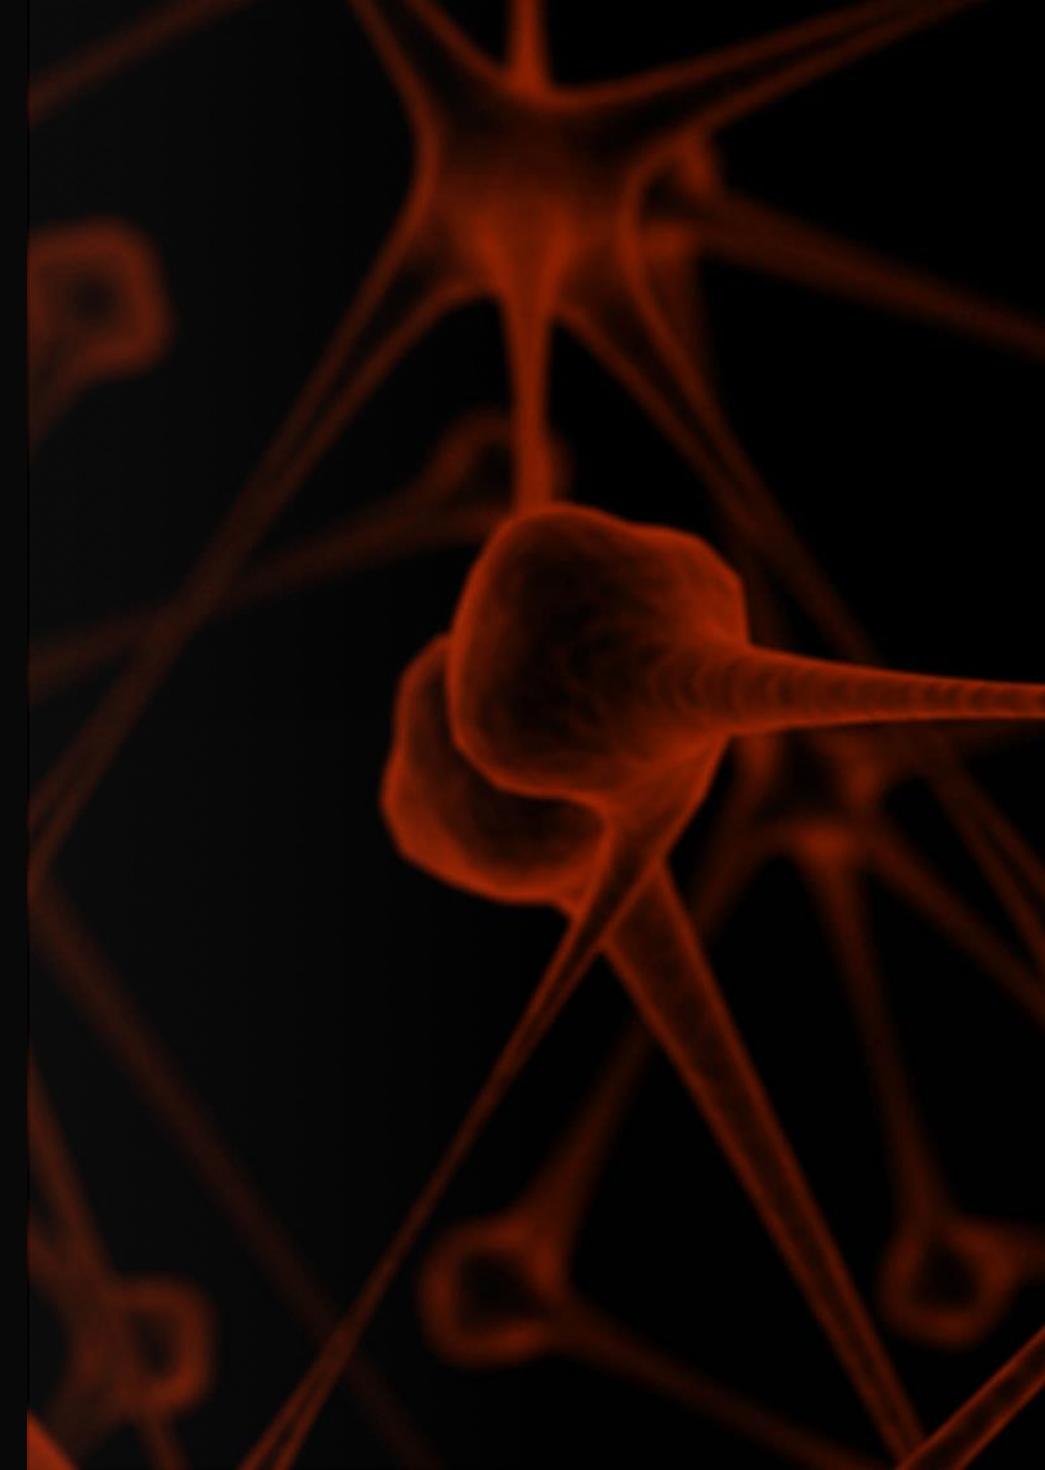
# Execute the Tensor Graph
with tf.Graph().as_default() as g:
    # The placeholders. Note the special name that end with '_input':
    y_input = tfs.block(df3, 'y', tf_name="y_input")
    z_input = tfs.block(df3, 'z', tf_name="z_input")

    # Perform elementwise sum and minimum
    y = tf.reduce_sum(y_input, [0], name='y')
    z = tf.reduce_min(z_input, [0], name='z')

    # The resulting dataframe
    (data_sum, data_min) = tfs.reduce_blocks([y, z], df3)

> # The final results are numpy arrays:
  print "Elementwise sum: %s and minimum: %s" % (data_sum, data_min)  
  
Elementwise sum: [ 45. -45.] and minimum: [ 0. -9.]
```

Tutorial Setup



Tutorial Prerequisites

- Log into Databricks Community Edition
(<https://community.cloud.databricks.com>)
- Import notebooks: <https://aka.ms/pydatatfs>
- Start cluster
- Install libraries
- Let's go!

Import Notebooks

Steps to Import notebooks

Repeat

Import notebook

The screenshot shows the Databricks web interface at the URL <https://community.cloud.databricks.com/?o=57901#>. The left sidebar has icons for Home, Workspace (selected), Recent, Data, and Clusters. The main workspace shows a list of notebooks under 'Shared' and 'PyData Seattle 2017'. A context menu is open over a notebook in the 'PyData Seattle 2017' folder, with 'Import' highlighted in blue.

Secure | https://community.cloud.databricks.com/?o=57901#

Workspace Shared PyData Seattle 2017

Documentation flights

Release Notes genomics

Training & Tutorials PyData Seattle 2017

Shared quick start

Users

Create

Clone

Rename

Move

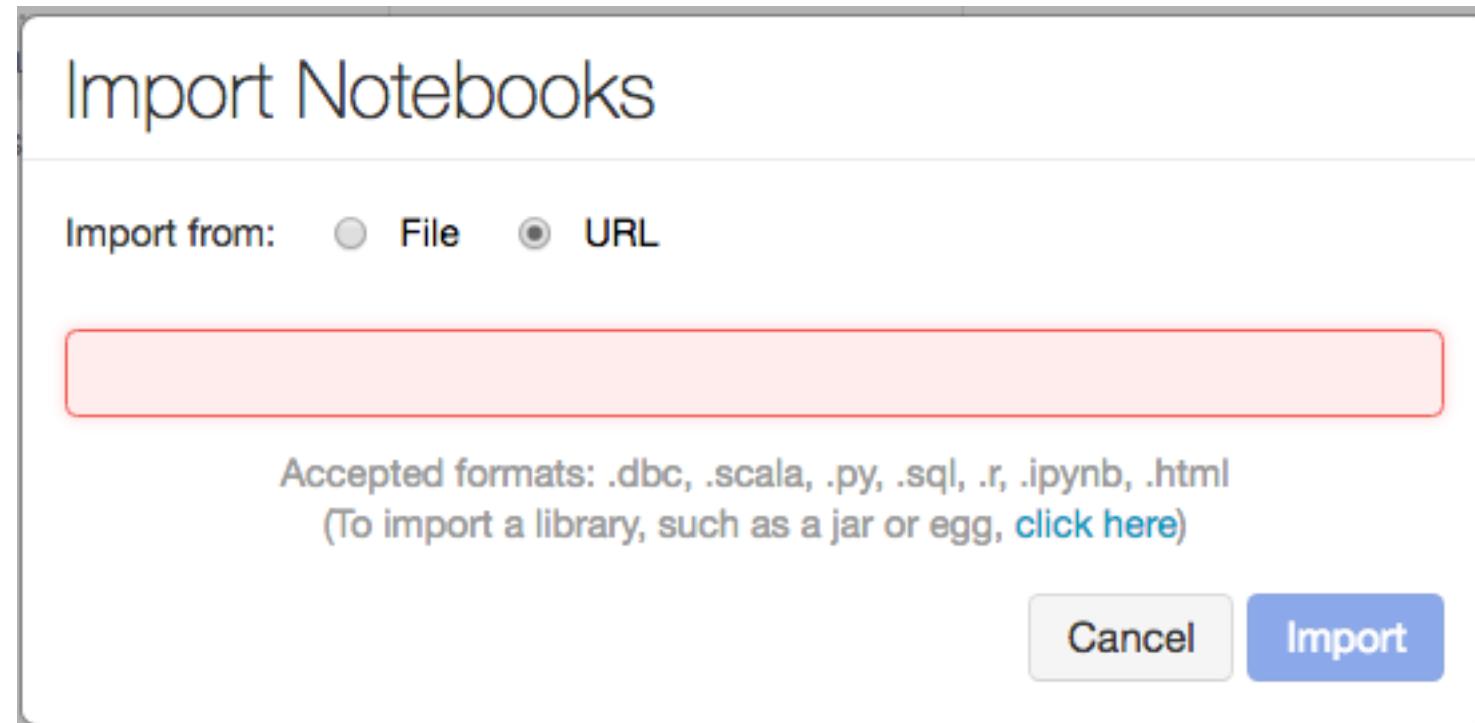
Delete

Import

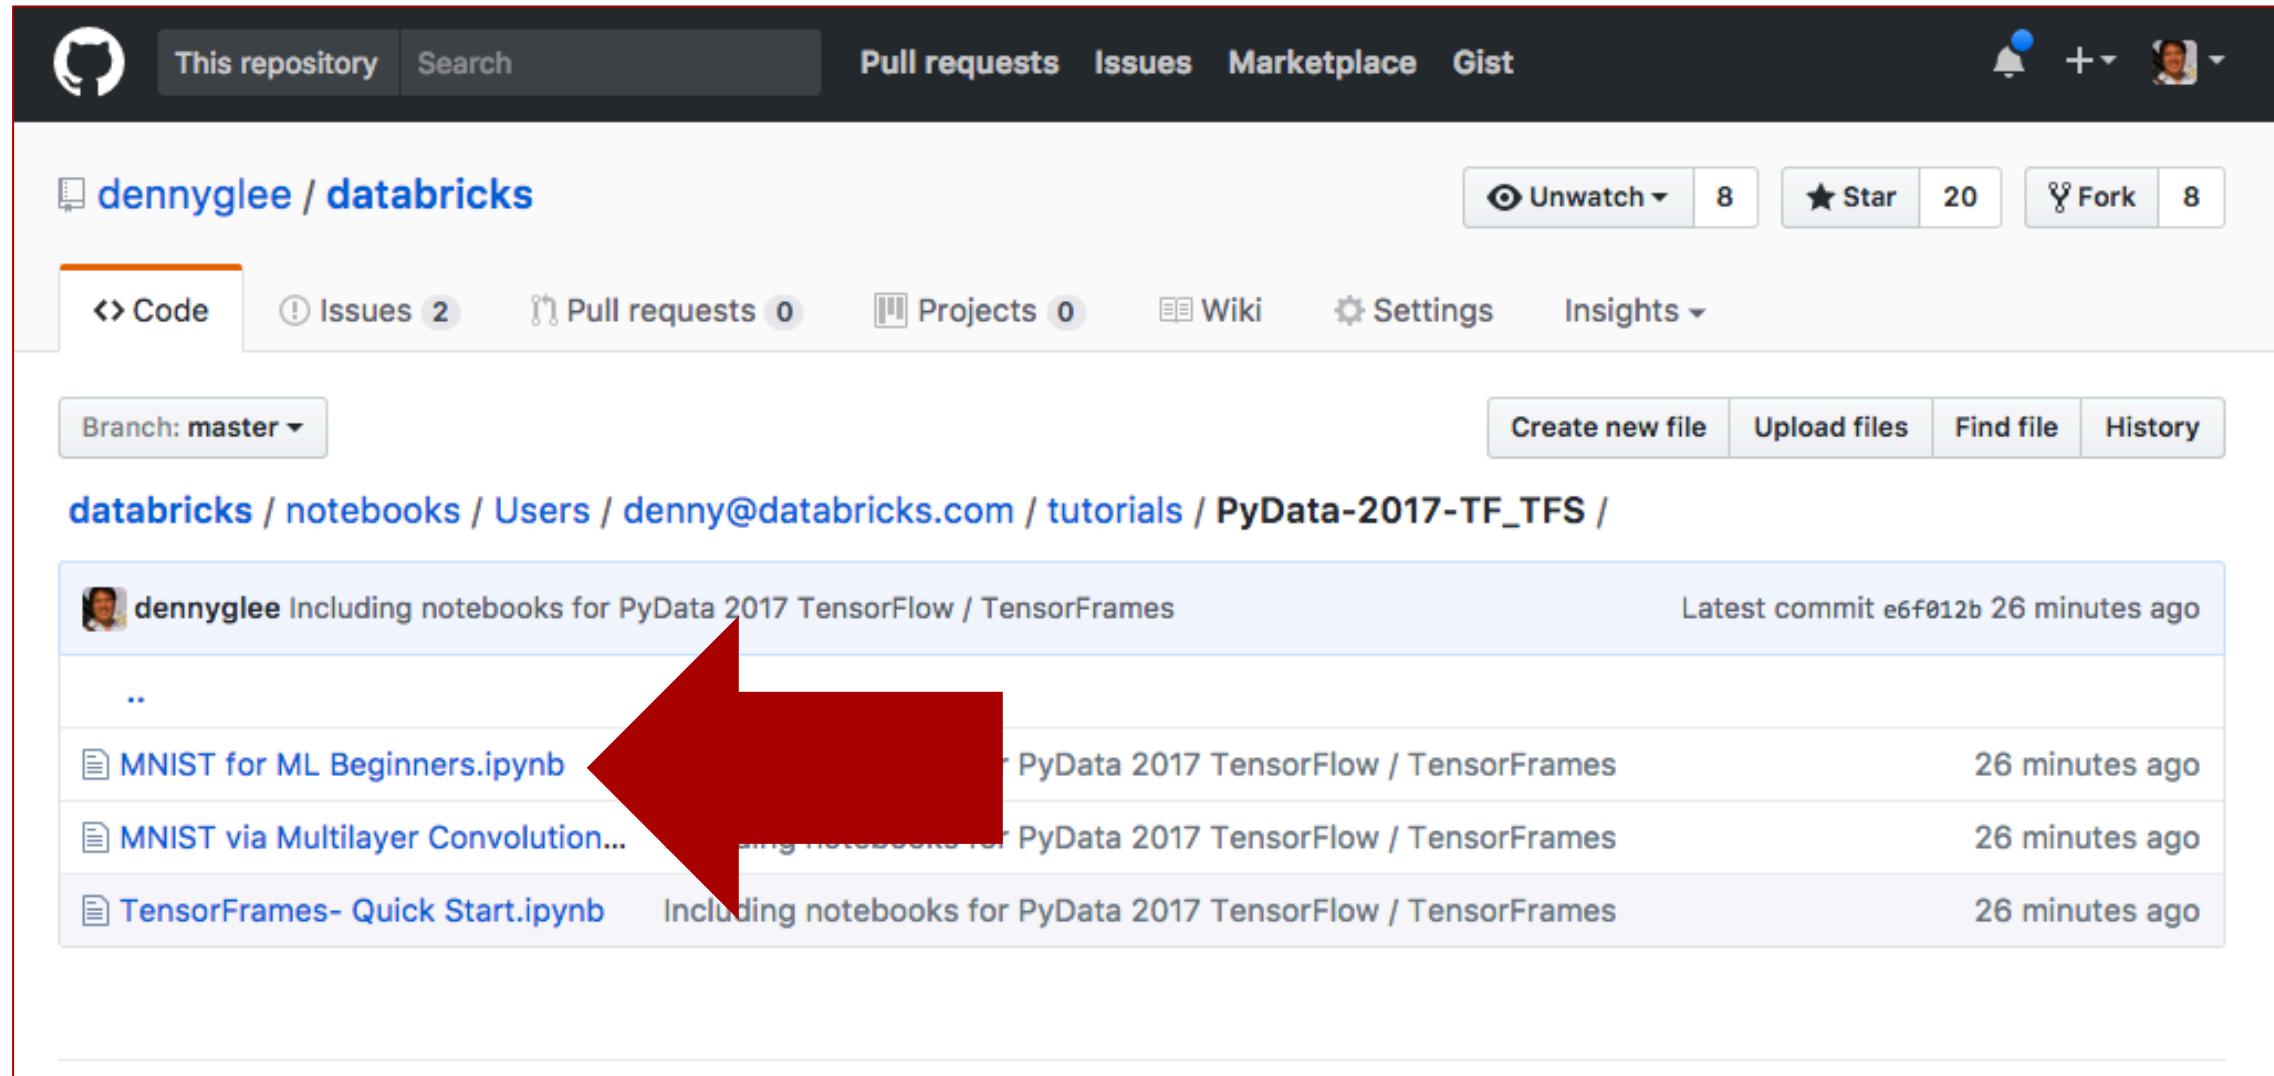
Export

Permissions

Import Notebook Dialog



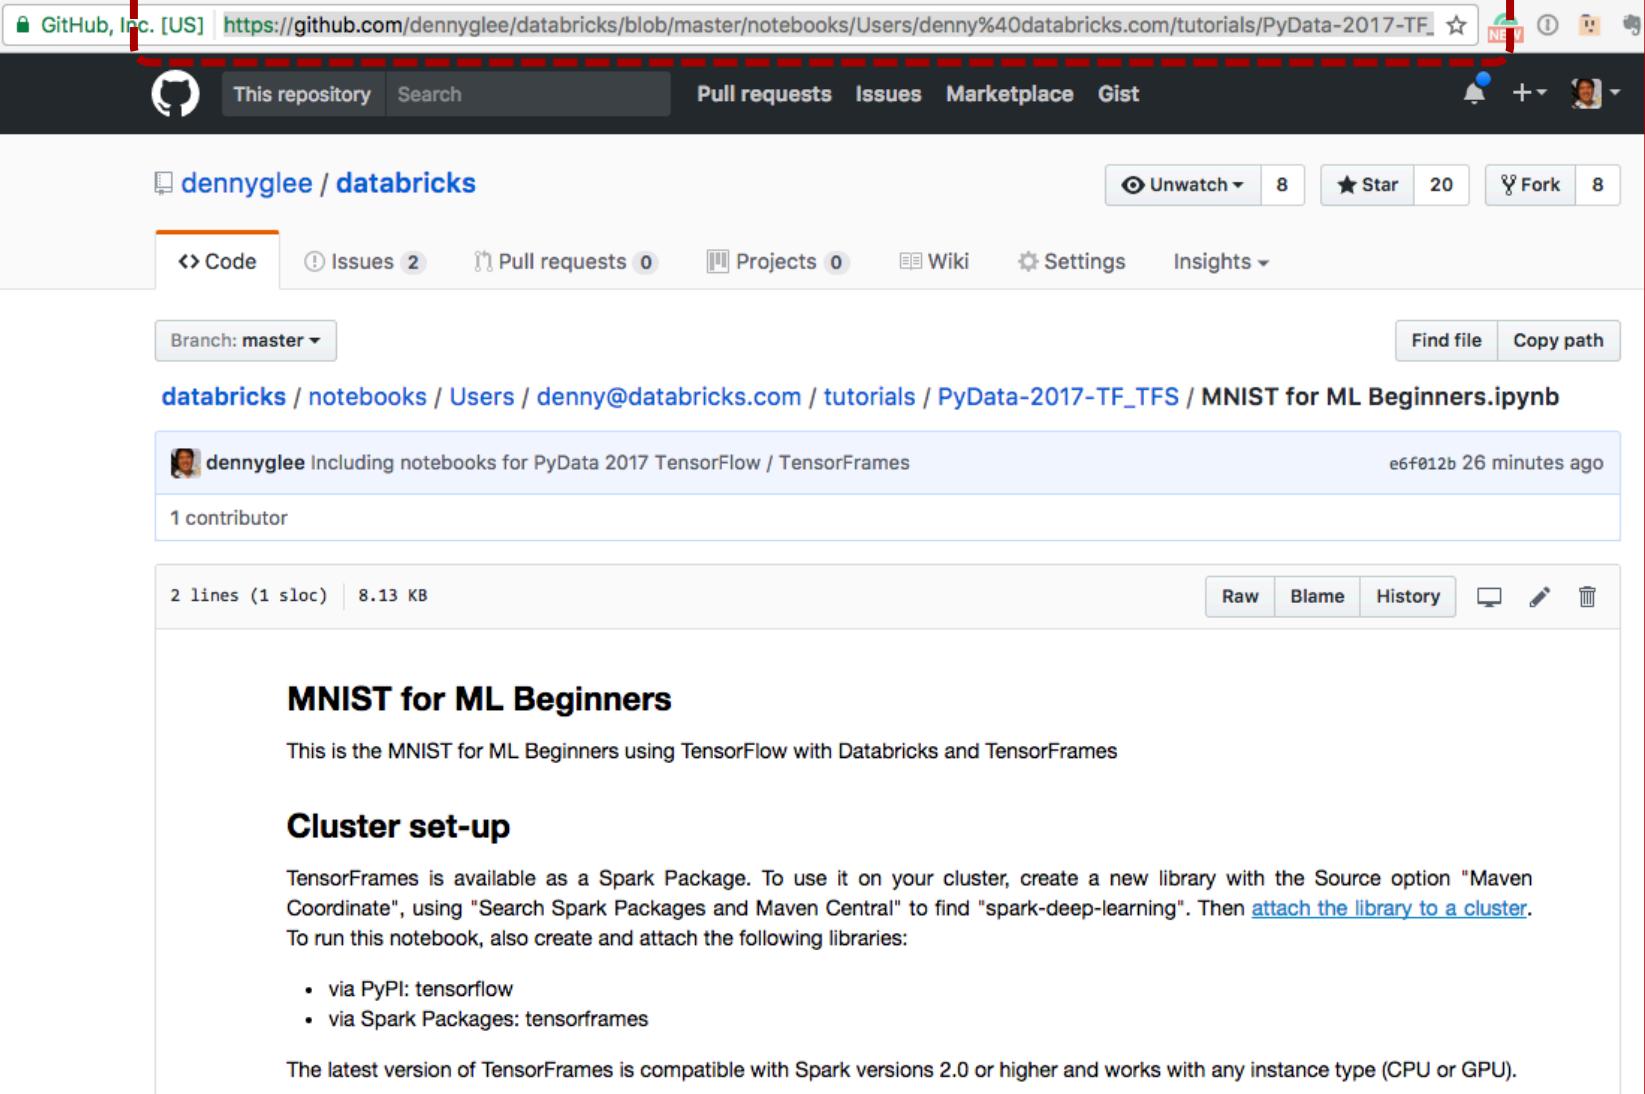
Goto <https://aka.ms/pydata-sea-2017-tfs>



A screenshot of a GitHub repository page for [dennyglee / databricks](#). The repository has 8 unwatched issues, 20 stars, and 8 forks. The Code tab is selected. The URL in the address bar is [databricks / notebooks / Users / denny@databricks.com / tutorials / PyData-2017-TF_TFS /](#). The repository contains several notebooks:

- dennyglee Including notebooks for PyData 2017 TensorFlow / TensorFrames Latest commit e6f012b 26 minutes ago
- ..
- MNIST for ML Beginners.ipynb Including notebooks for PyData 2017 TensorFlow / TensorFrames 26 minutes ago
- MNIST via Multilayer Convolutional.ipynb Including notebooks for PyData 2017 TensorFlow / TensorFrames 26 minutes ago
- TensorFrames- Quick Start.ipynb Including notebooks for PyData 2017 TensorFlow / TensorFrames 26 minutes ago

Open a notebook > Get URL



The screenshot shows a GitHub repository page for the user 'dennylee' with the repository name 'databricks'. The specific notebook 'MNIST for ML Beginners.ipynb' is highlighted. The browser's address bar at the top contains the URL https://github.com/dennylee/databricks/blob/master/notebooks/Users/denny%40databricks.com/tutorials/PyData-2017-TF_TFS/MNIST%20for%20ML%20Beginners.ipynb.

MNIST for ML Beginners
This is the MNIST for ML Beginners using TensorFlow with Databricks and TensorFrames

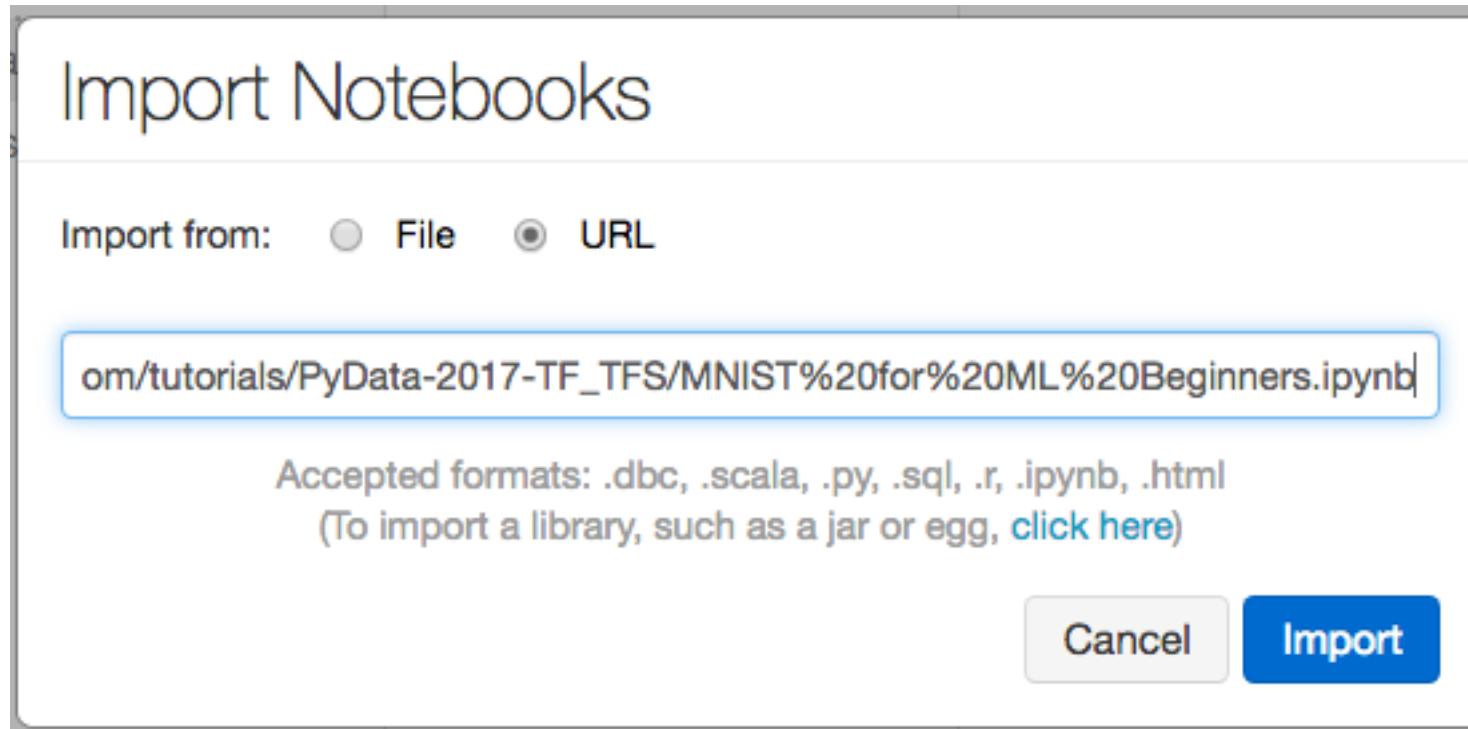
Cluster set-up

TensorFrames is available as a Spark Package. To use it on your cluster, create a new library with the Source option "Maven Coordinate", using "Search Spark Packages and Maven Central" to find "spark-deep-learning". Then [attach the library to a cluster](#). To run this notebook, also create and attach the following libraries:

- via PyPI: tensorflow
- via Spark Packages: tensorframes

The latest version of TensorFrames is compatible with Spark versions 2.0 or higher and works with any instance type (CPU or GPU).

Paste URL into Import Notebooks dialog



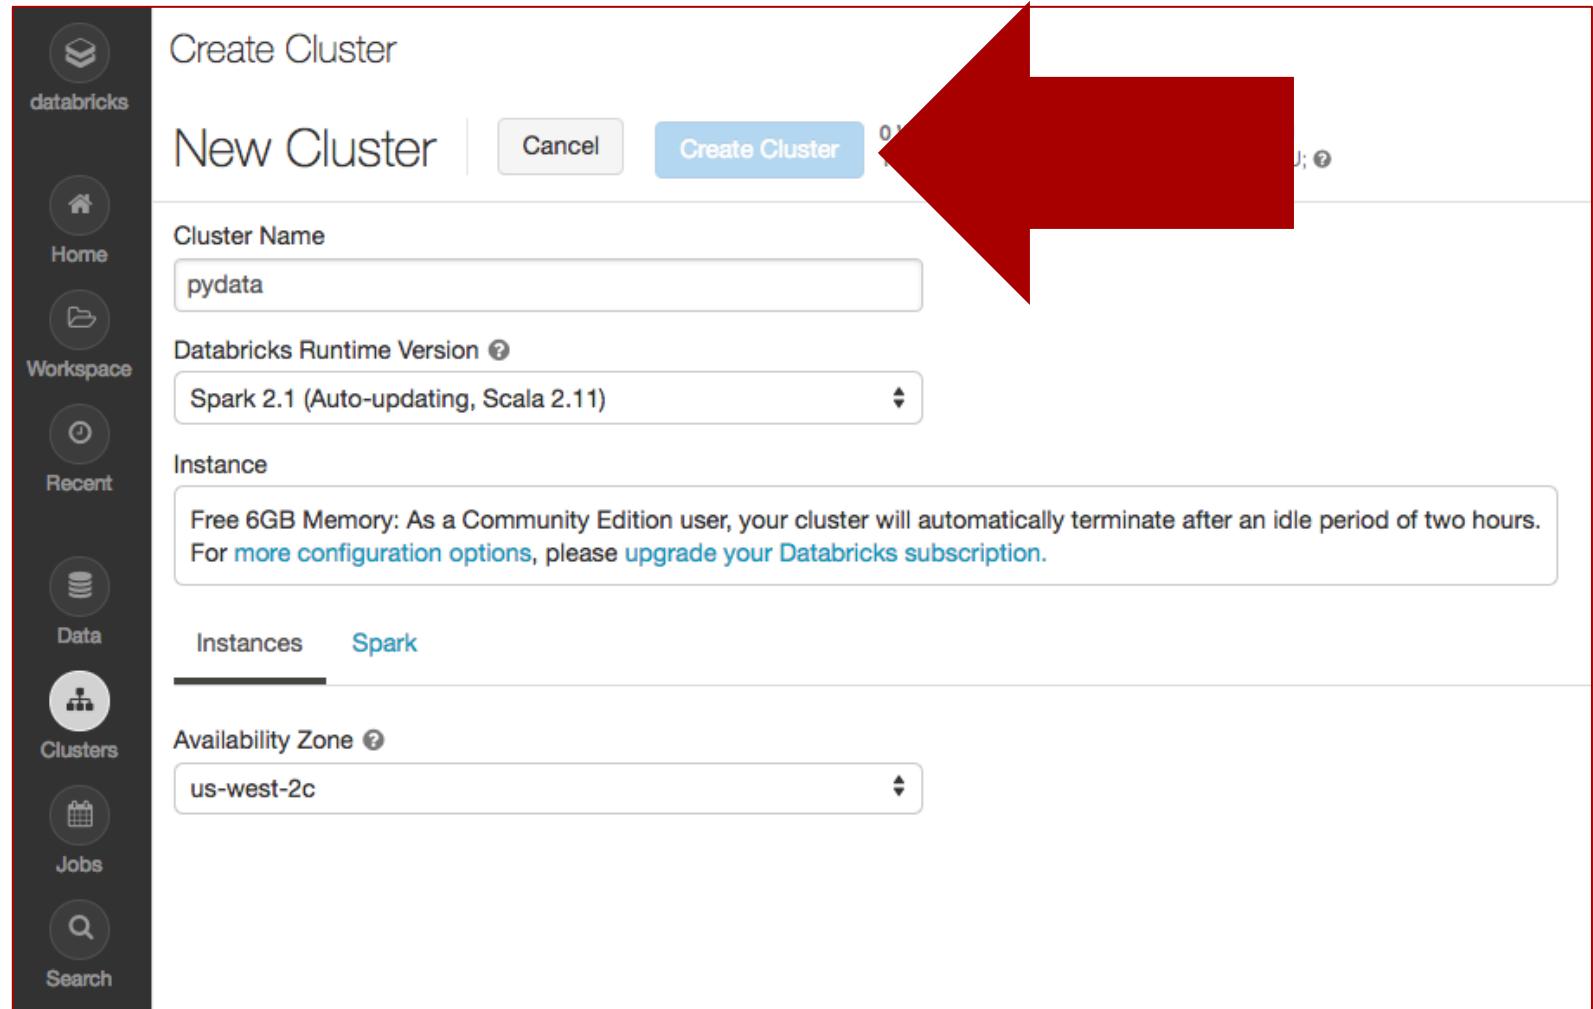
Do it for all three notebooks

Start a cluster

Three steps to start a Spark cluster

Start a cluster

1. Clusters > Create Cluster
2. Fill in name and Run-time version: Spark 2.1 (Scala 2.11)



Install libraries

Installing TensorFlow via PyPi

Installing TensorFrames via Spark Packages

Create Library

The screenshot shows the Databricks workspace interface. On the left is a dark sidebar with icons for Home, Workspace, Recent, Data, Clusters, Jobs, and Search. The main area has a top navigation bar with 'Workspace' and 'Users' dropdowns. Under 'Users', 'denny.g.lee@gmail.com' is selected. A context menu is open over the 'Create' button, with 'Library' highlighted. Below the menu, there's a section for 'Featured Notebooks'.

Workspace

- Documentation
- Release Notes
- Training & Tutorials
- Shared
- Users

Users

- denny.g.lee@gmail.com
- deborah.siegel@gmail.com
- drabas.t@gmail.com
- ali@databricks-demo....
- matei@databricks-de...

Create

- Library
- Folder

Featured Notebooks

- Notebook
- Library
- Folder

Notebook

Job

Cluster

Table

Library

Install TensorFlow via PyPi

New Library

Language

Upload Python Egg or PyPI



Install PyPi Package

You can specify a package name with an optional [version specification](#)

PyPi Name

tensorflow

Install Library

Attach TensorFlow to your cluster

The screenshot shows the Databricks interface. On the left is a dark sidebar with icons for Databricks, Home, Workspace, Recent, Data, and Clusters. The 'Clusters' icon is highlighted with a red border. The main area has a red border and contains the following content:

tensorflow

tensorflow | x Delete

PyPI rules

[tensorflow](#)

Clusters

Attach automatically to all clusters.

Attach	Name	Status
<input checked="" type="checkbox"/>	pandas	Attached

Install TensorFrames via Spark Packages

New Library

Source **Maven Coordinate**

Install Maven Artifacts

Coordinate **Maven Coordinate (e.g. com.databricks:spark-csv_2.10:1.0.0)**

Search Spark Packages and Maven Central

▶ Advanced Options

Create Library

Find TensorFrames

Search Packages

tensorframes 

Spark Packages 

Name	Organization	Description	Rating	Releases	Options
tensorframes	databricks	Tensorflow wrapper for DataFrames on Apache Spark	5 / 5	0.2.8-s_2.11 	+ Select
tensorframes	tjhunter	Tensorflow wrapper for DataFrames on Apache Spark	0 / 5	0.2.2-s_2.10 	+ Select



Close

Create TensorFrames Library ... Almost!

New Library

Source

Maven Coordinate



Install Maven Artifacts

Coordinate

databricks:tensorframes:0.2.8-s_2.11

Search Spark Packages and Maven Central

► Advanced Options

Create Library

Attach TensorFrames Library

tensorframes-0.2.8-s_2.11

tensorframes-0.2.8-s_2.11

 Delete

Artifacts

[scalactic_2.11-3.0.0.jar](#)
[commons-proxy-1.0.jar](#)
[commons-lang3-3.4.jar](#)
[scala-logging-slf4j_2.11-2.1.2.jar](#)
[scala-logging-api_2.11-2.1.2.jar](#)
[libtensorflow-1.1.0-rc1.jar](#)
[slf4j-api-1.7.7.jar](#)
[tensorflow-1.1.0-rc1.jar](#)
[tensorframes-0.2.8-s_2.11.jar](#)
[libtensorflow_jni-1.1.0-rc1.jar](#)

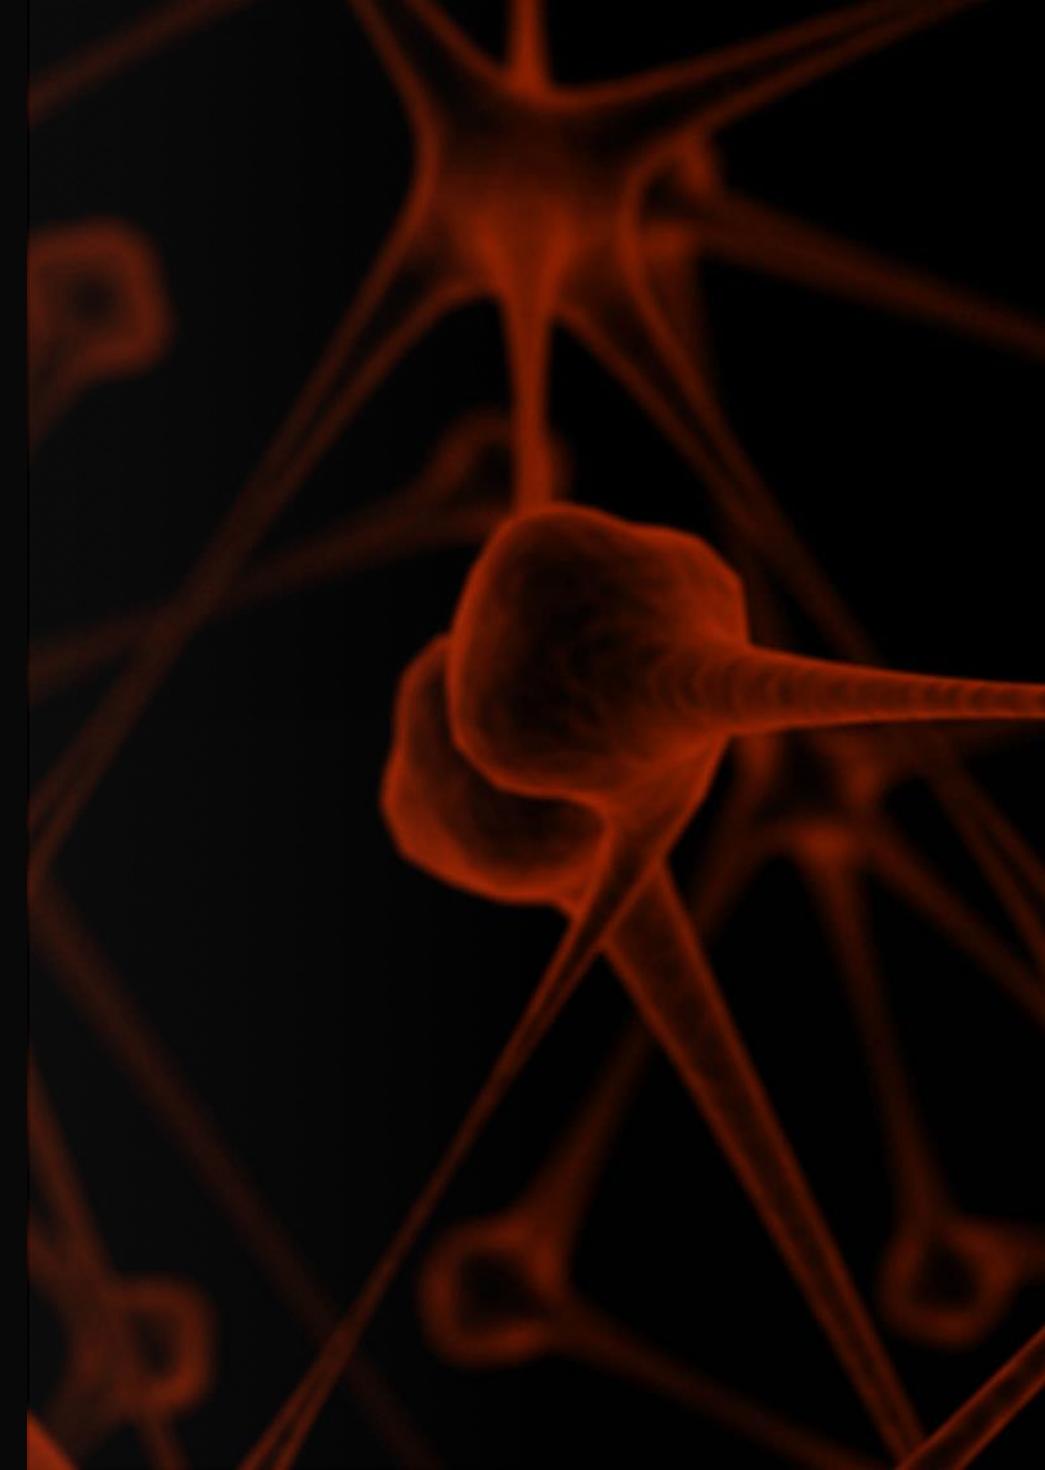
Clusters

Attach automatically to all clusters.

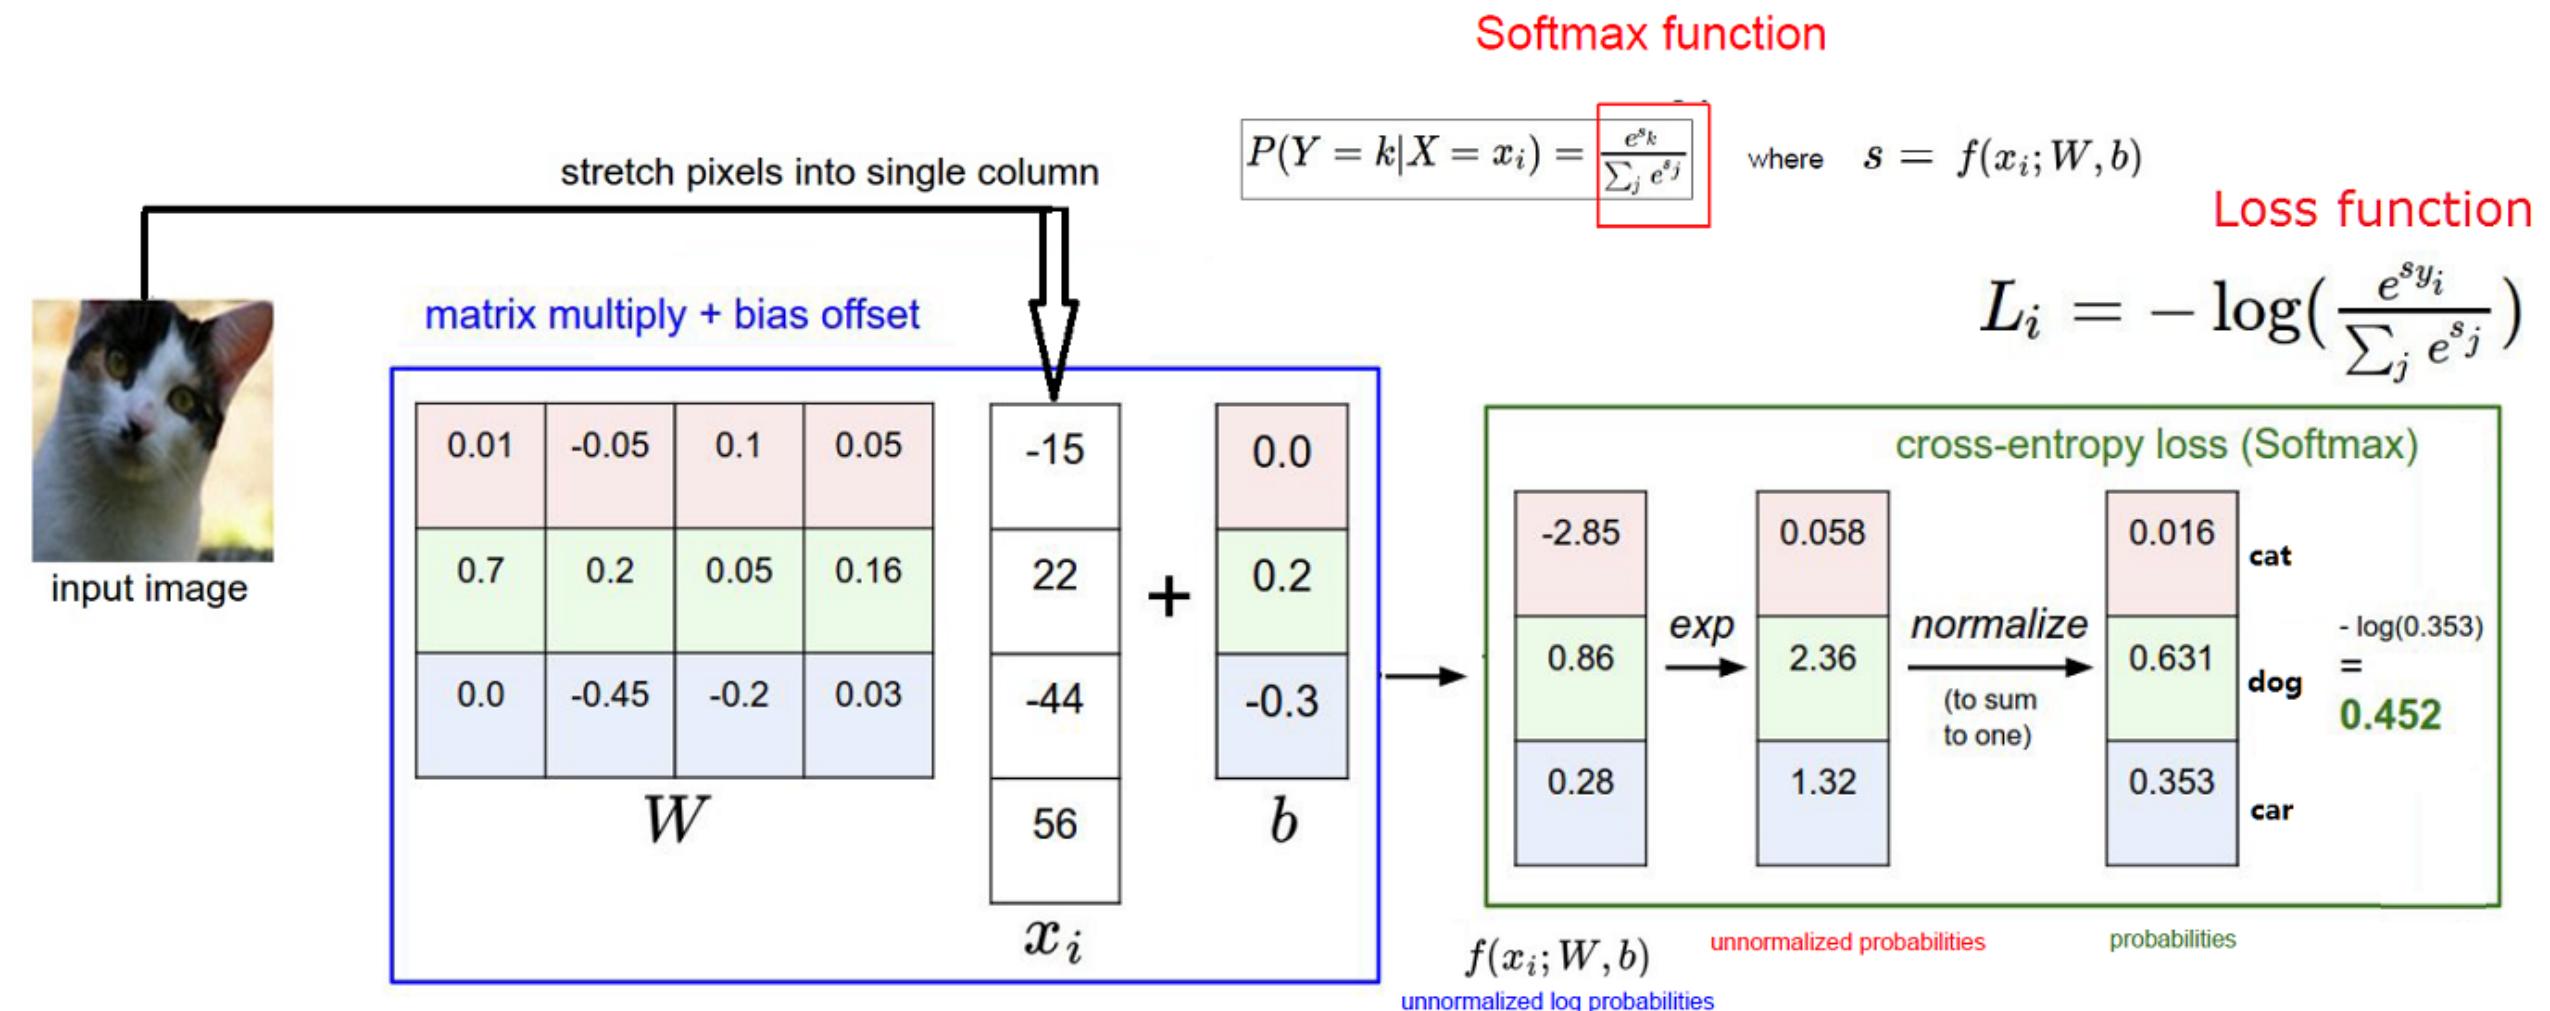
Attach	Name	Status
<input checked="" type="checkbox"/>	pandas	Attached

Let's go!

Handwritten Digit Recognition using Softmax



Softmax Function (Linear Classifier)



Source: <http://bit.ly/2tNRZpg> | Demo: <http://stanford.io/2tJCtd7>

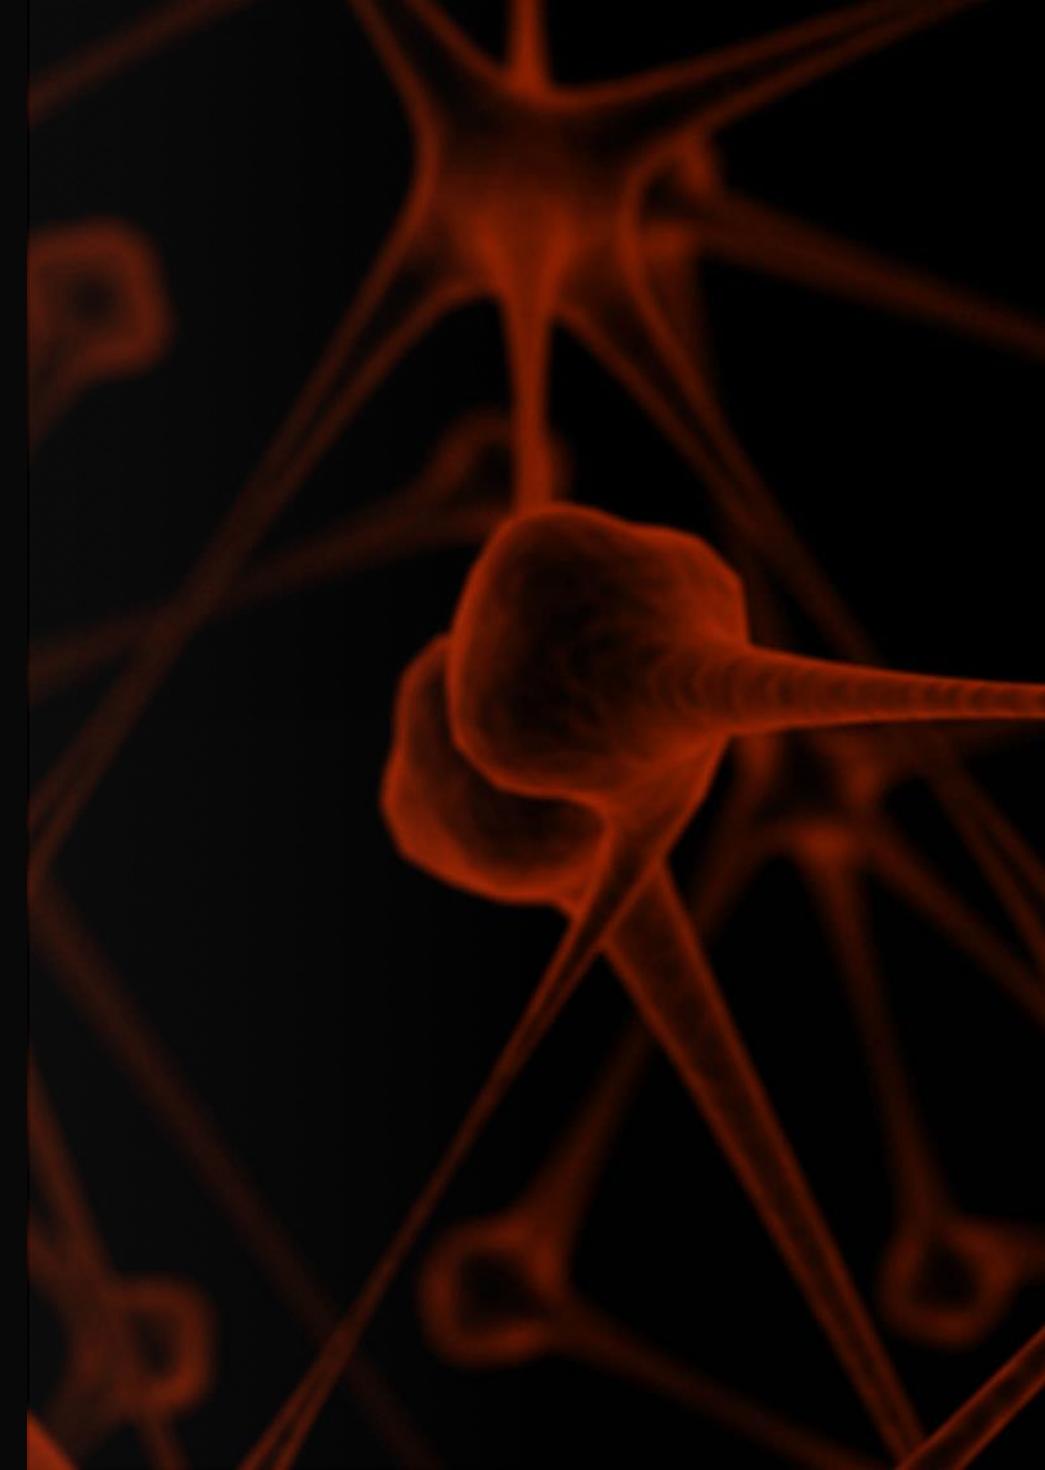
CNN accuracy: ~92

```
1 correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
2 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
3 print(accuracy.eval(feed_dict={x: mnist.test.images, y_: mnist.test.labels}))
4 #print(sess.run(accuracy, feed_dict={x: mnist.test.images, y_: mnist.test.labels}))
```

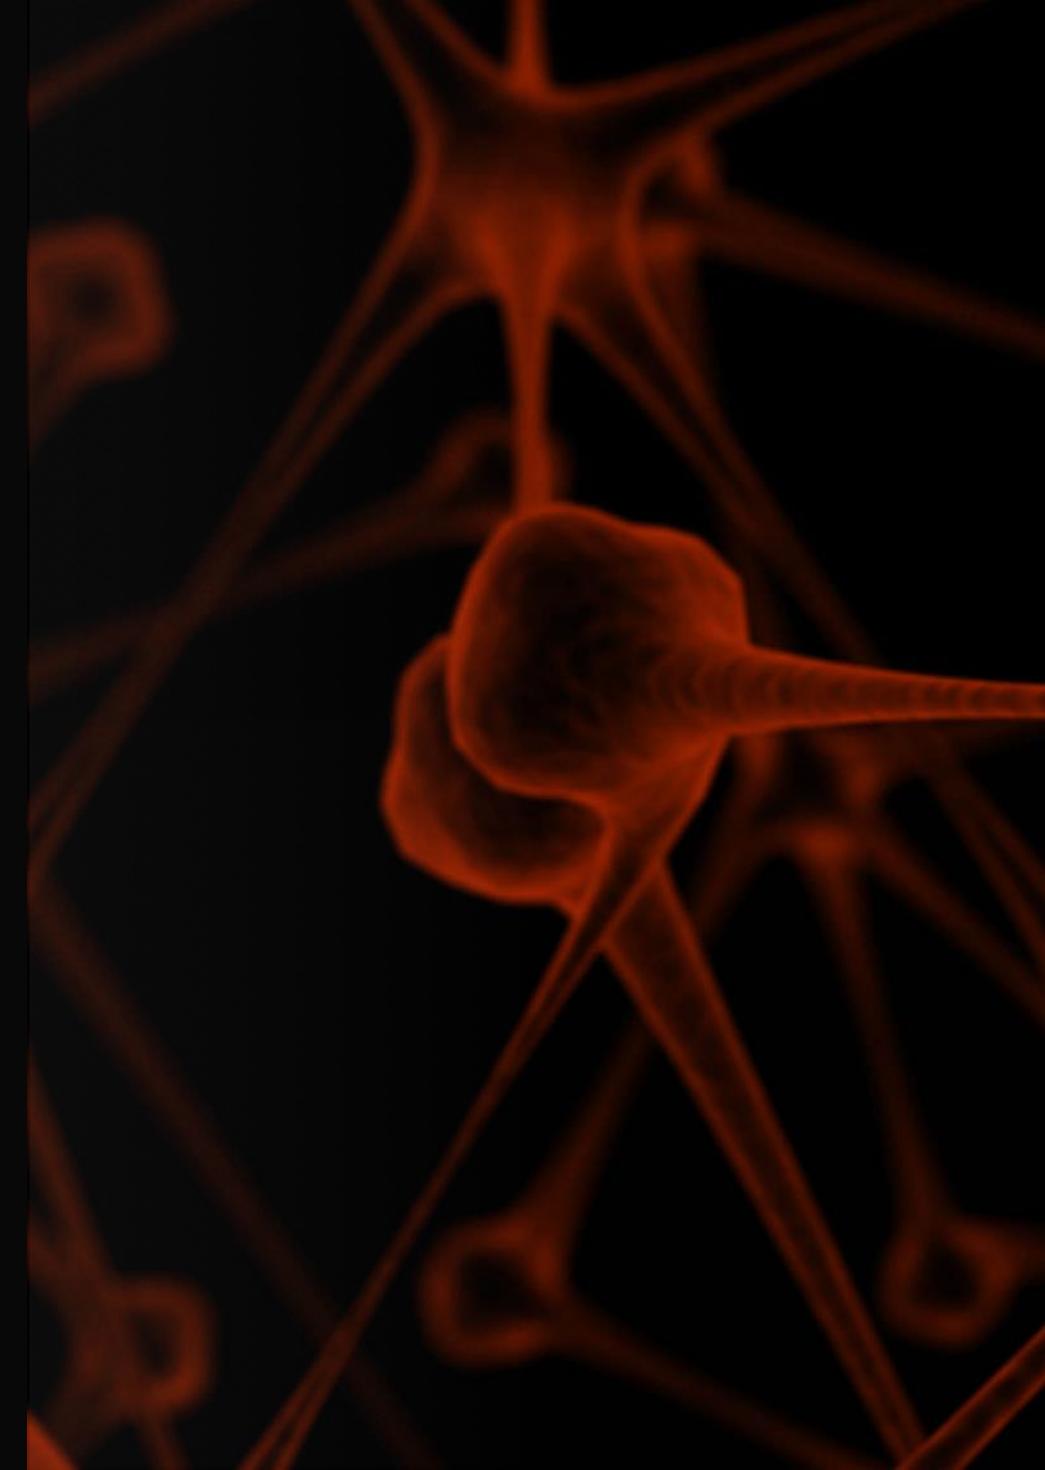
0.9071

ConvNetJS MNIST Demo

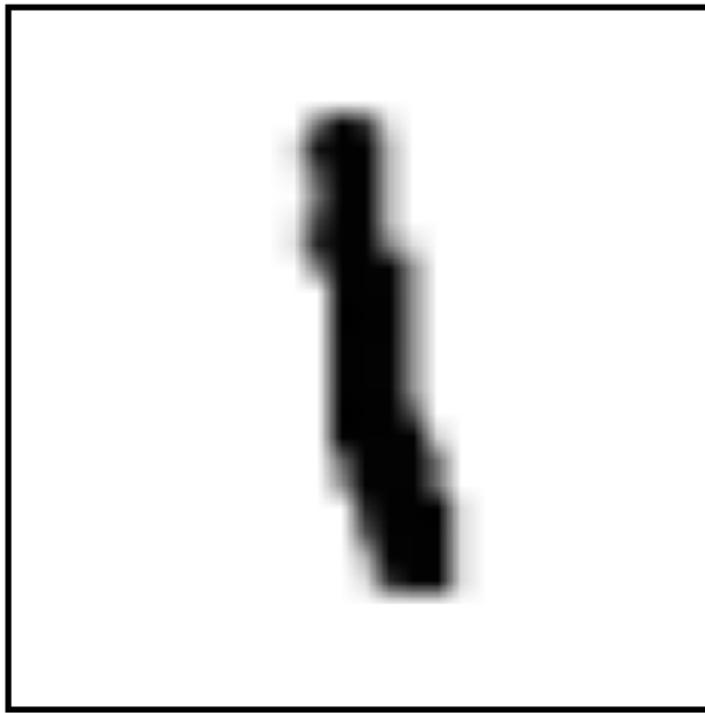
Andrej Karpathy's MNIST Demo:
[https://cs.stanford.edu/people/karpathy/
convnetjs/demo/mnist.html](https://cs.stanford.edu/people/karpathy/convnetjs/demo/mnist.html)



Handwritten Digit Recognition using CNNs



Handwritten Digit Recognition



2

Source: <http://bit.ly/2sbFW0R>

MNIST Data

- Set of handwritten digits based on the National Institute of Science and Technology – i.e. a modified dataset or [MNIST dataset](#) – via LeCun et. al
- Original source is [NIST Special Database 19 Handprinted Forms and Characters Database](#) comprised of:
 - Images: Each digit is a scanned 28px x 28px flattened into an array of size 784
 - Labels: One-Hot vector representing the actual number associated with the image

Great References

[Andrej Karpathy's ConvNetJS MNIST Demo](#)

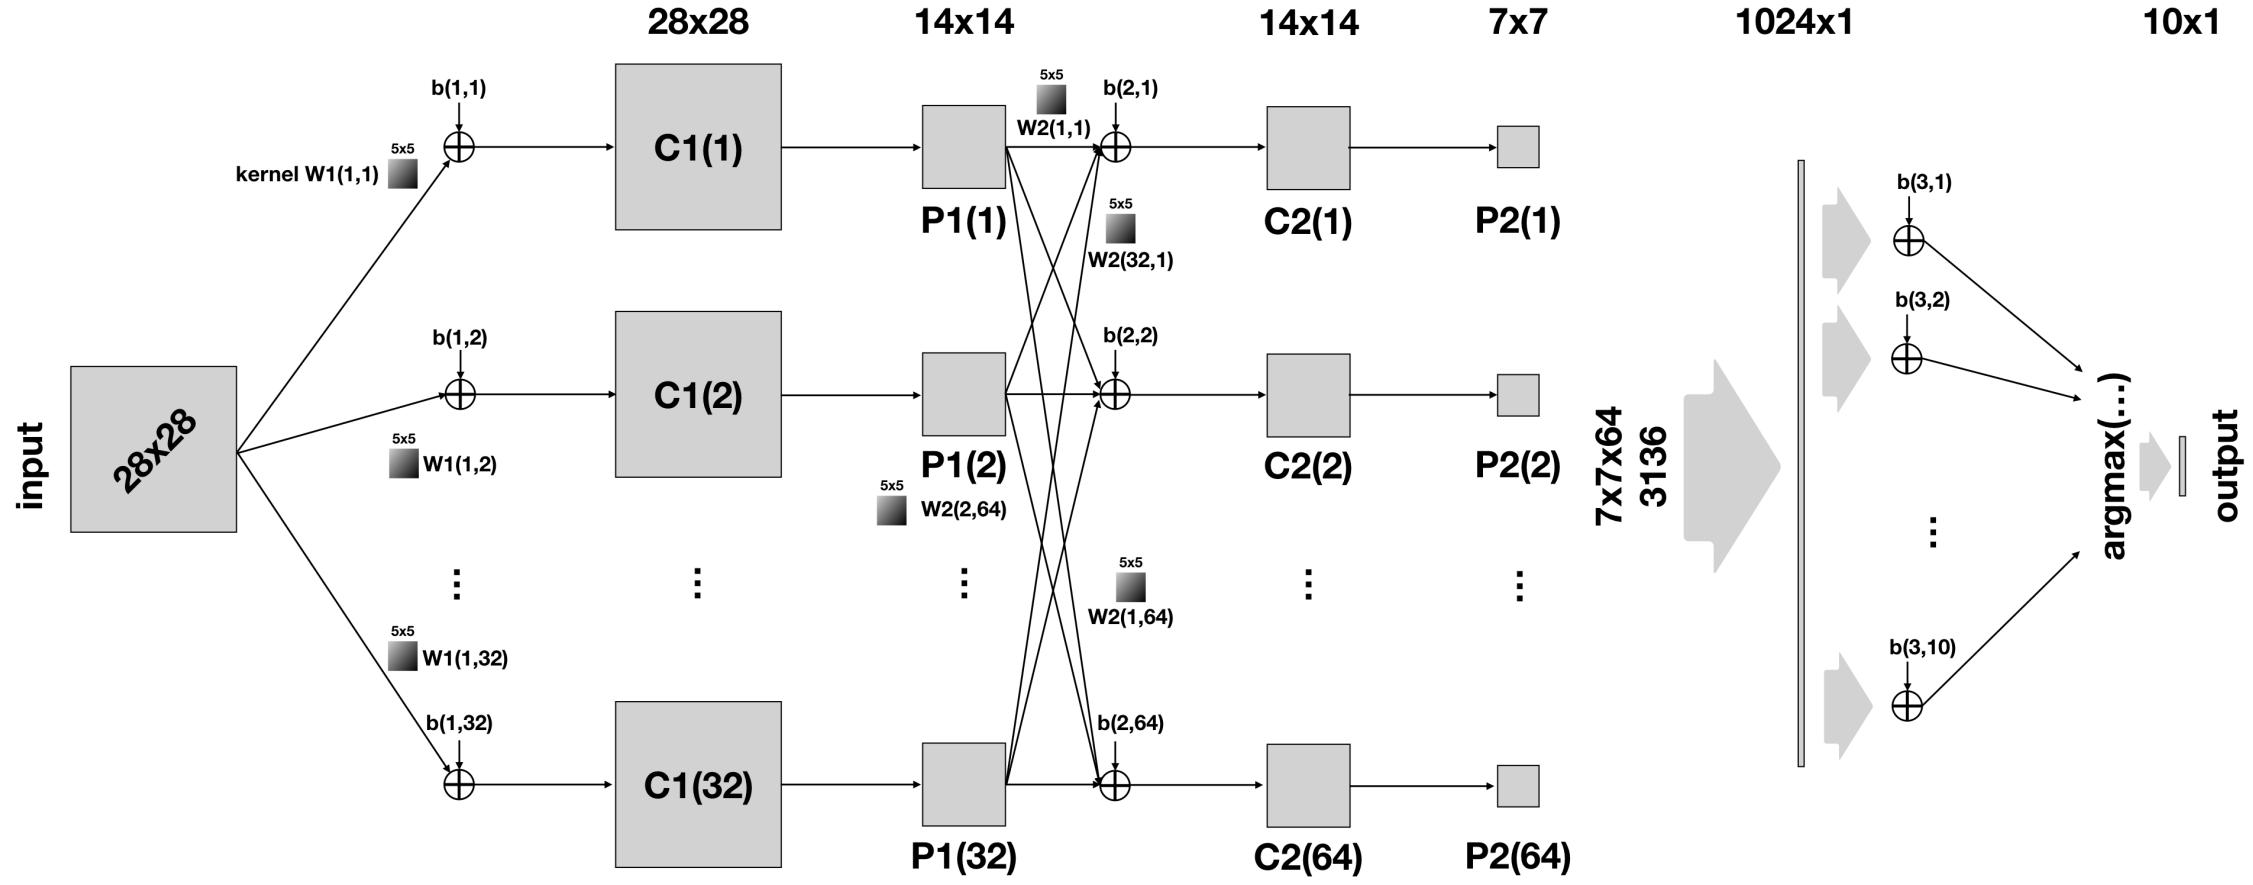
CS231n: Convolutional Neural Networks for Visual Recognition

[Syllabus and Slides](#) | [Course Notes](#) | [YouTube](#)

- With particular focus on [CS231n: Lecture 7: Convolution Neural Networks](#)

Upcoming webinar: [Build, Scale, and Deploy Deep Learning Pipelines with Ease](#)

CNN Structure



Convolution 1

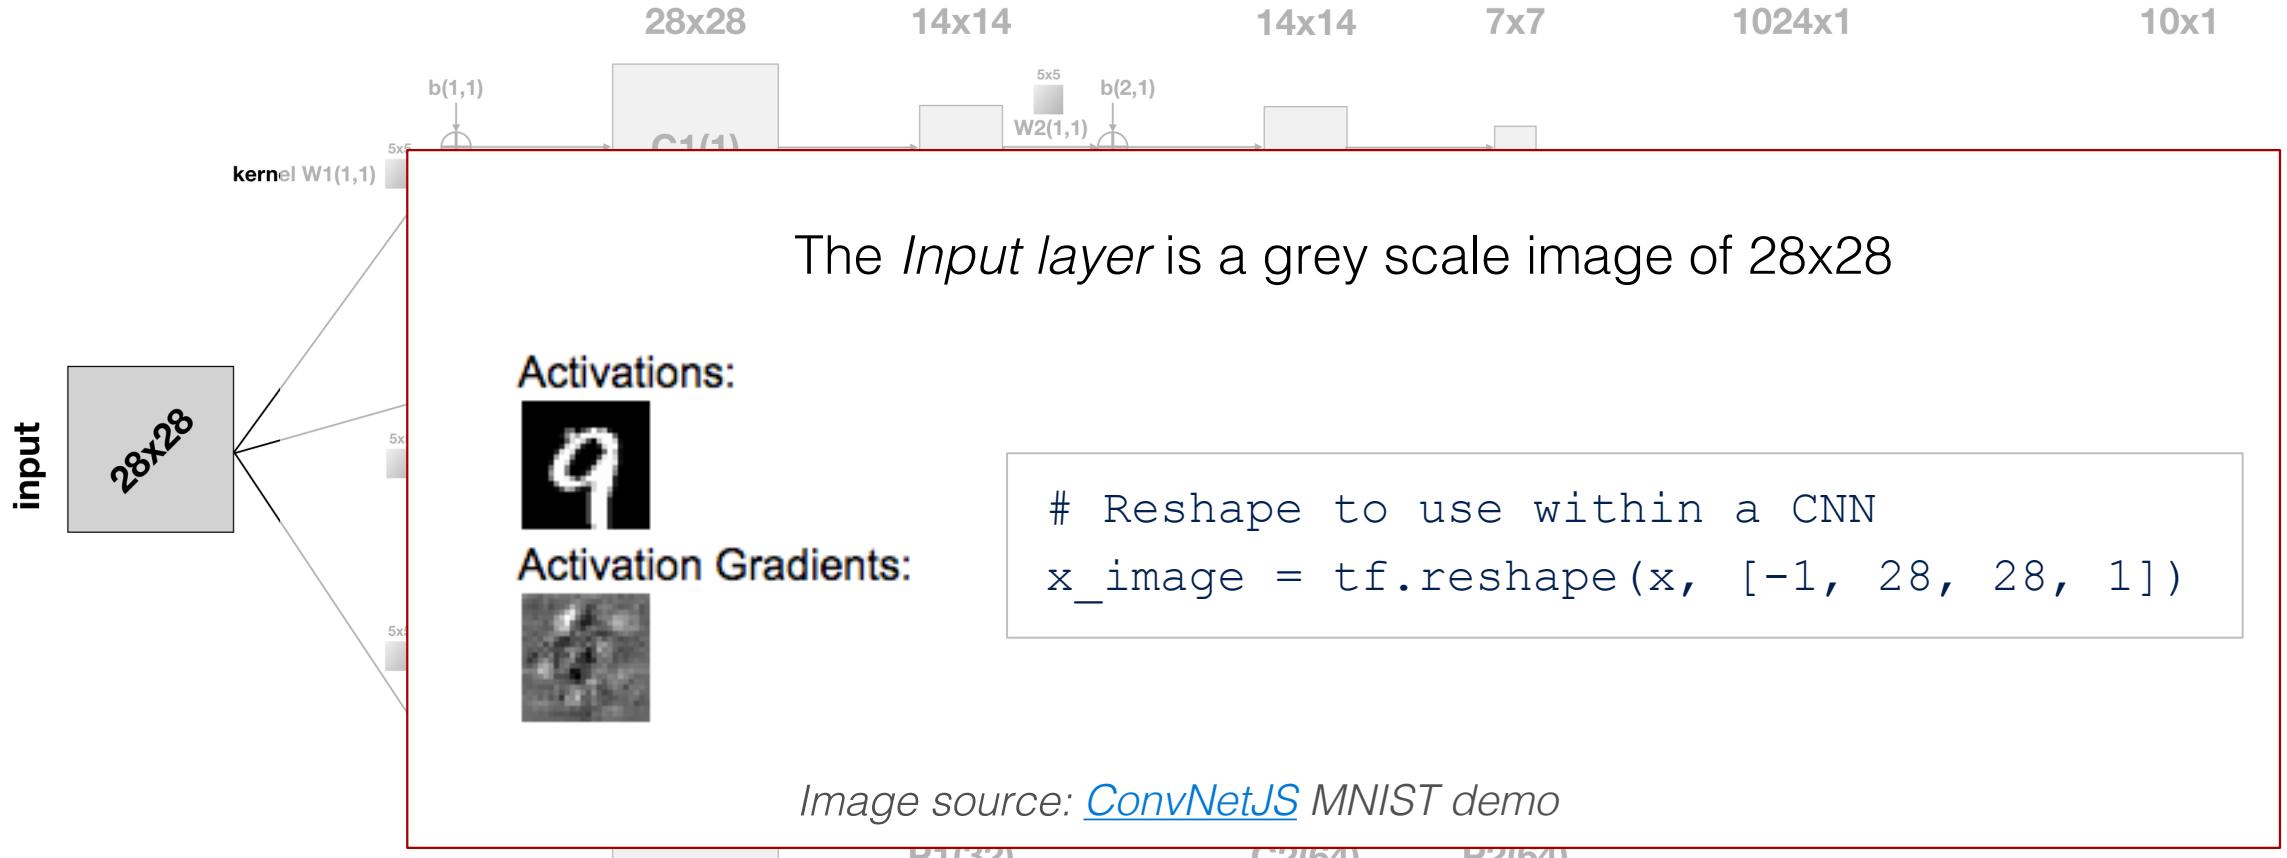
Pooling 1

Convolution 2

Pooling 2

Fully connected
feed-forward

CNN Structure



Convolution 1

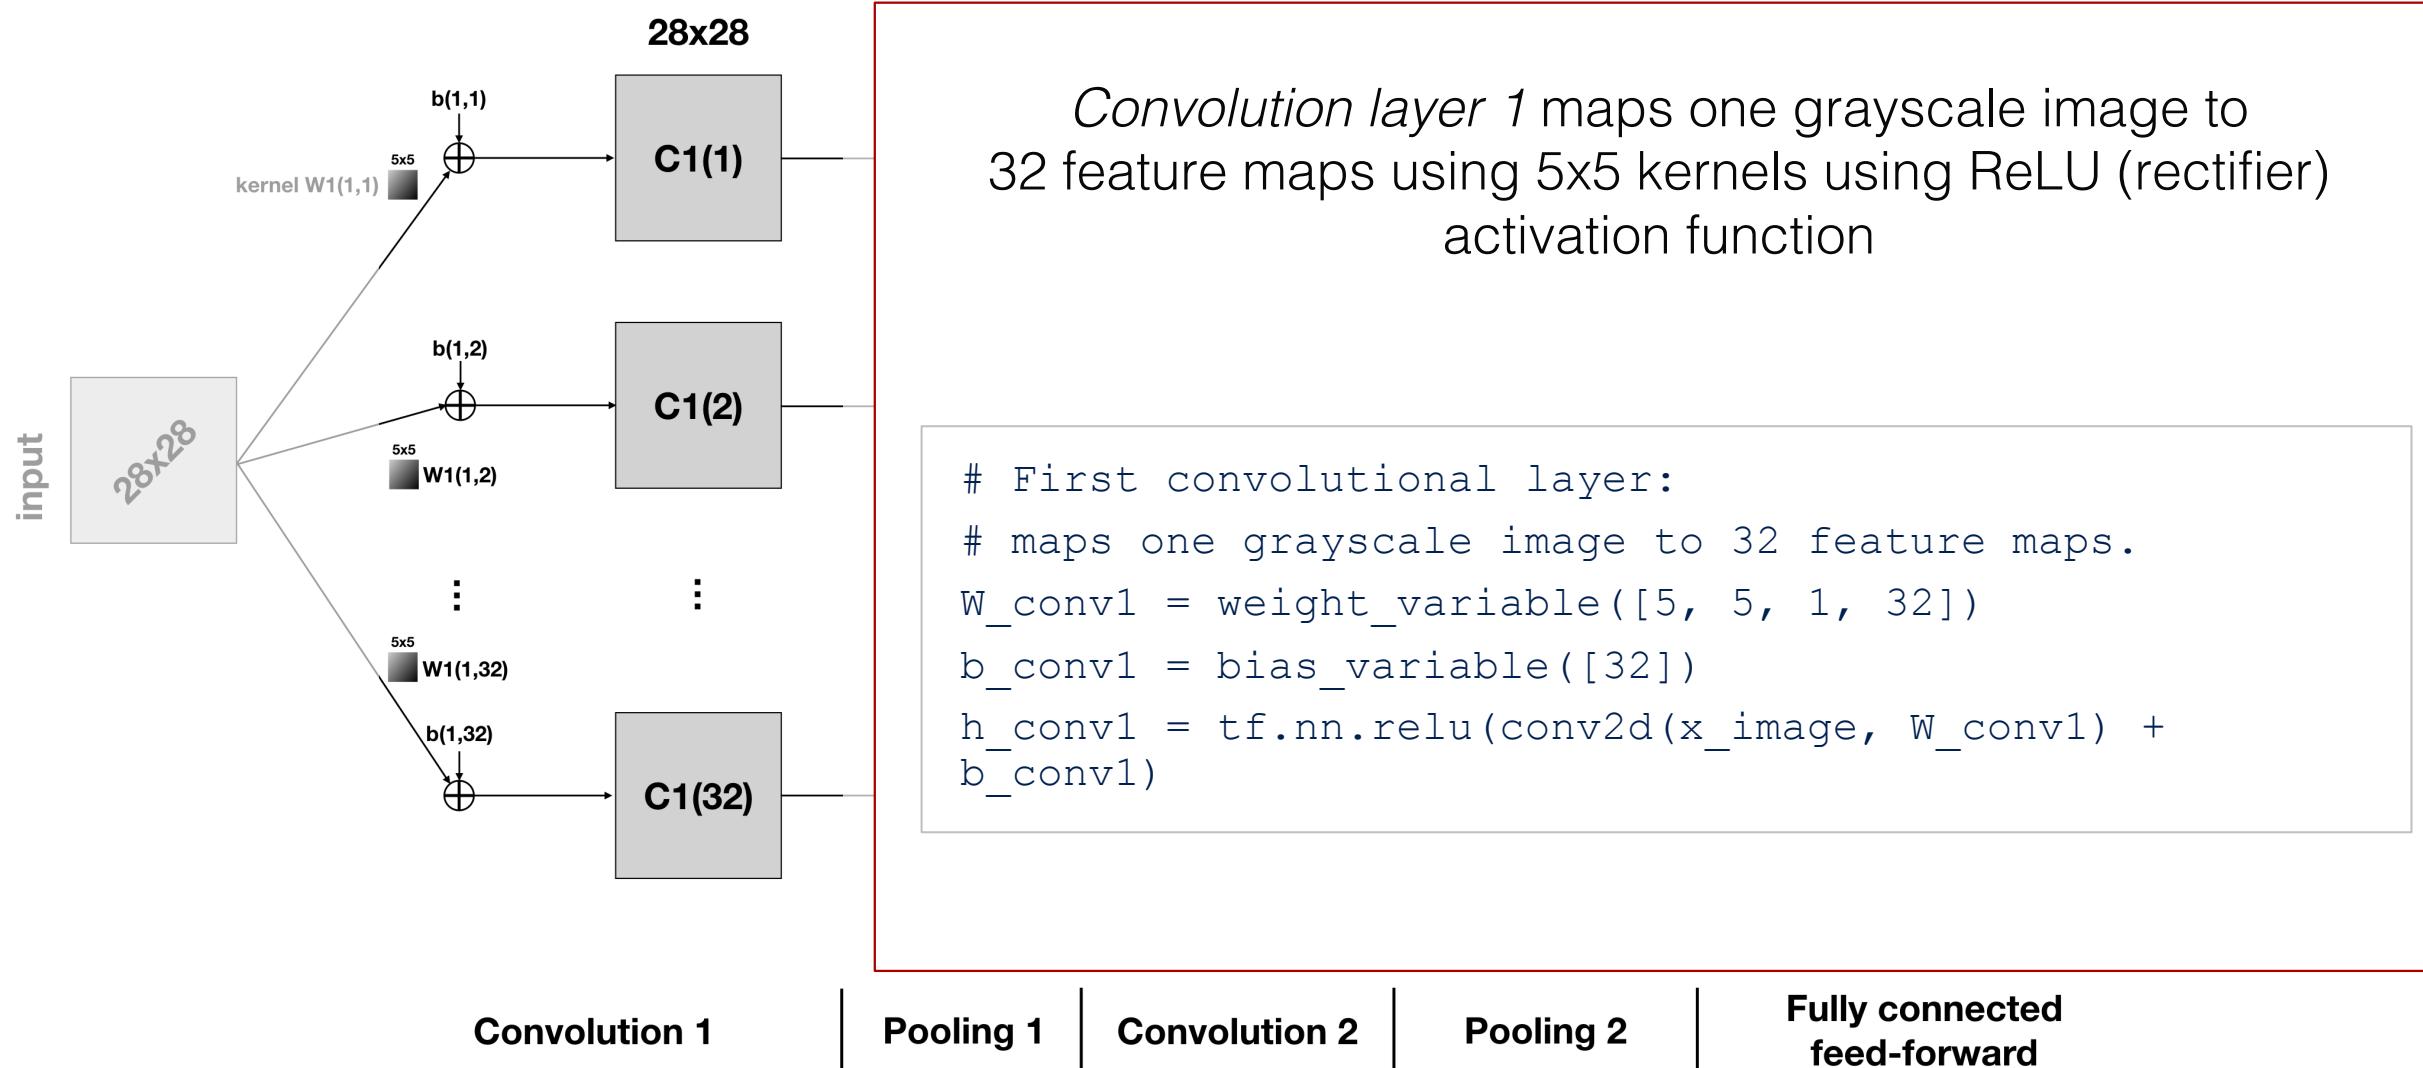
Pooling 1

Convolution 2

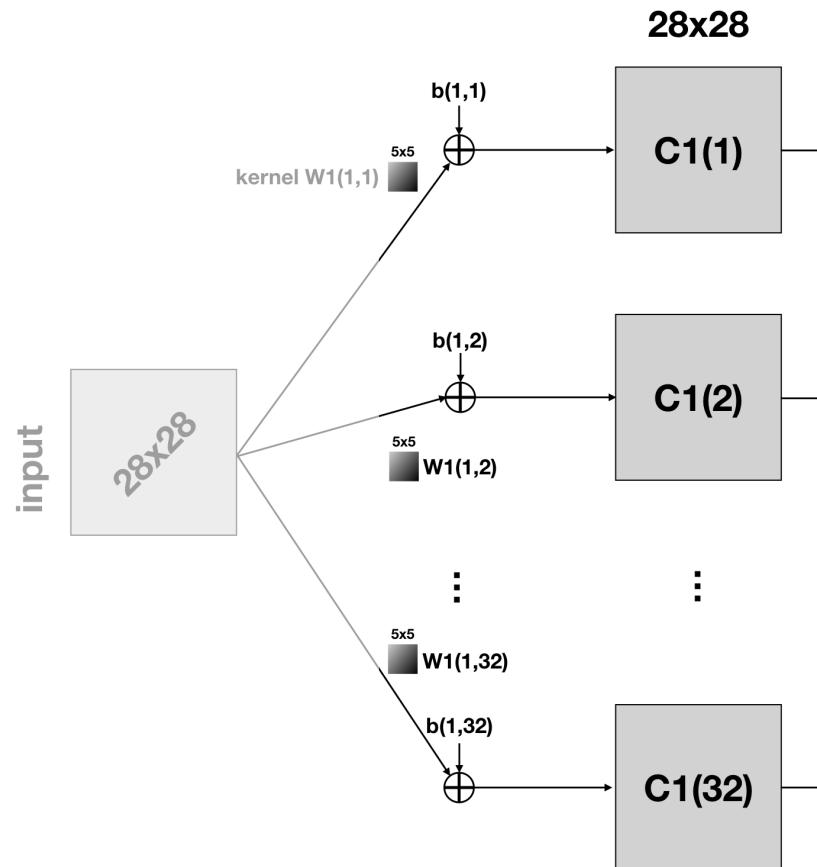
Pooling 2

Fully connected
feed-forward

CNN Structure: Convolution Layer 1



CNN Structure: Convolution Layer 1



Convolution layer 1 maps one grayscale image to 32 feature maps using 5×5 kernels using ReLU (rectifier) activation function

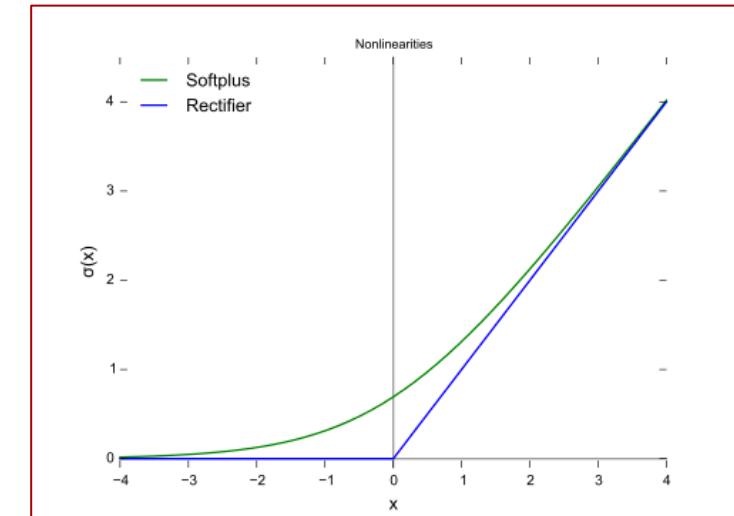
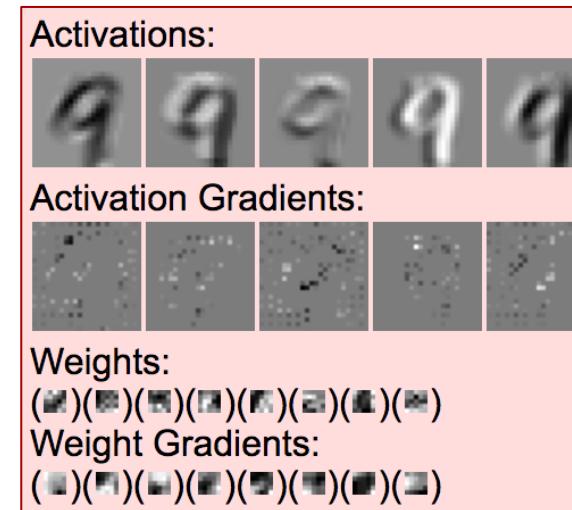


Image source: [ConvNetJS](#) MNIST demo

Convolution 1

Pooling 1

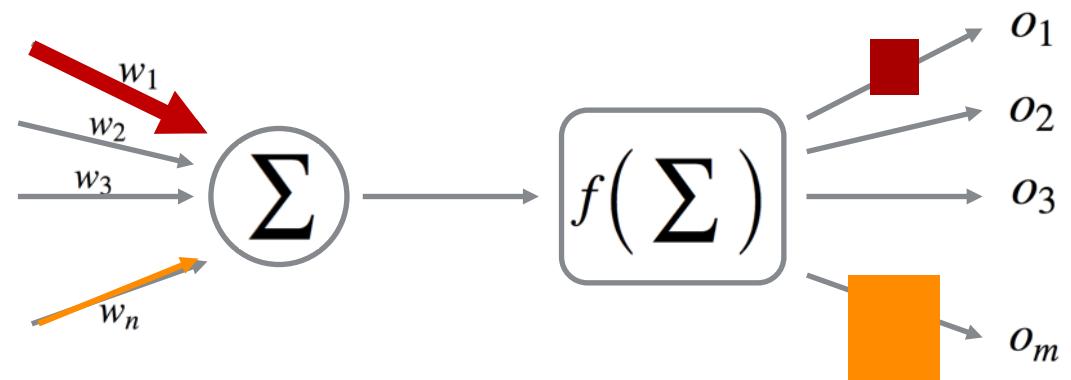
Convolution 2

Pooling 2

**Fully connected
feed-forward**

Why normalize? (continued)

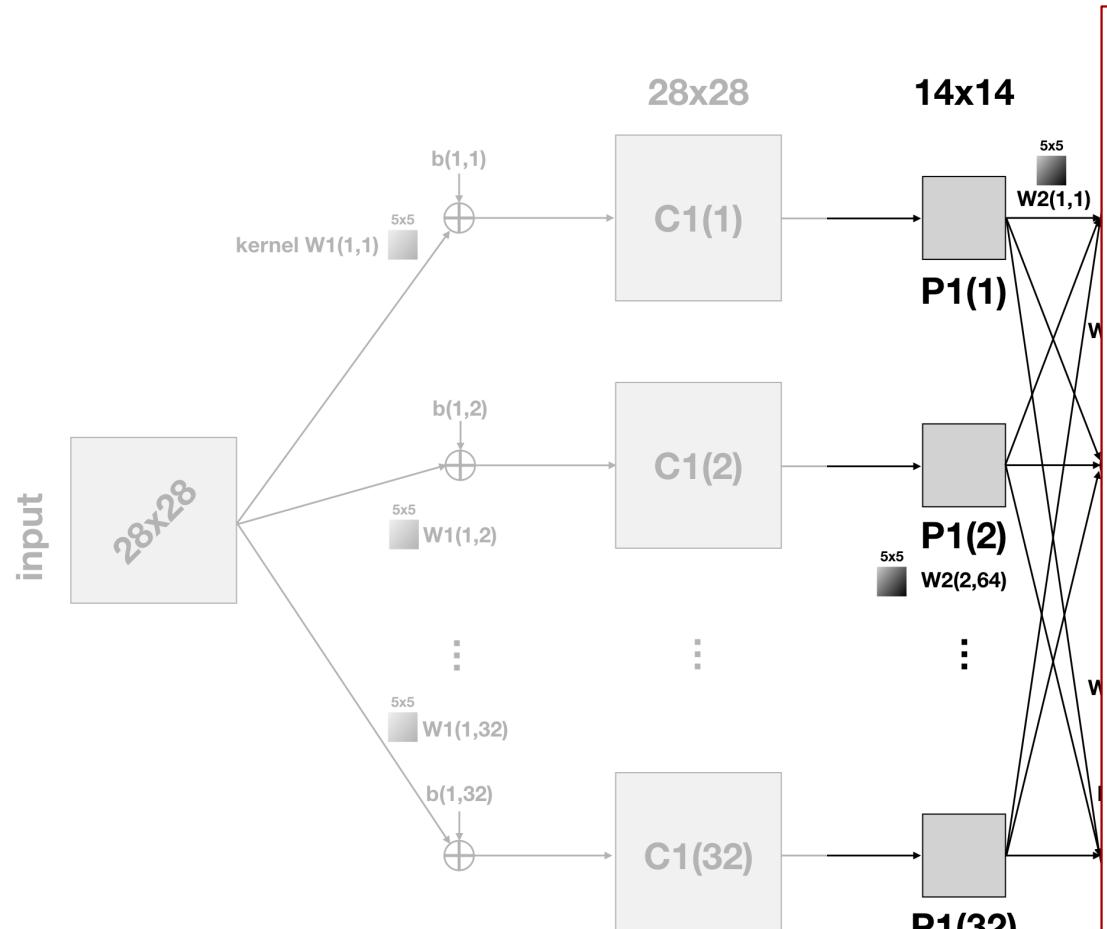
- Normalization changes the range of pixel intensity
- Also known as contrast stretching or histogram stretching
- Bring image values to the same range



“If we don’t scale our input training vectors, the ranges of distribution of those feature vectors would likely be different for each vector – i.e. learning rate would cause corrections, e.g. over compensating one and under compensating another”

[Source: <http://bit.ly/2rsc2b1>]

CNN Structure: Pooling 1



*Pooling layer 1 down samples image by 2x
so you have a 14x14 matrix*

```
# Pooling layer - downsamples by 2X.  
h_pool1 = max_pool_2x2(h_conv1)
```

Activations:



Activation Gradients:

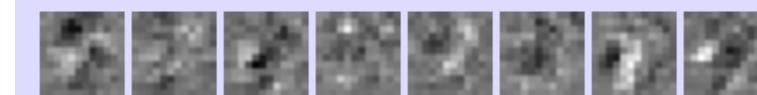


Image source: [ConvNetJS](#) MNIST demo

Convolution 1

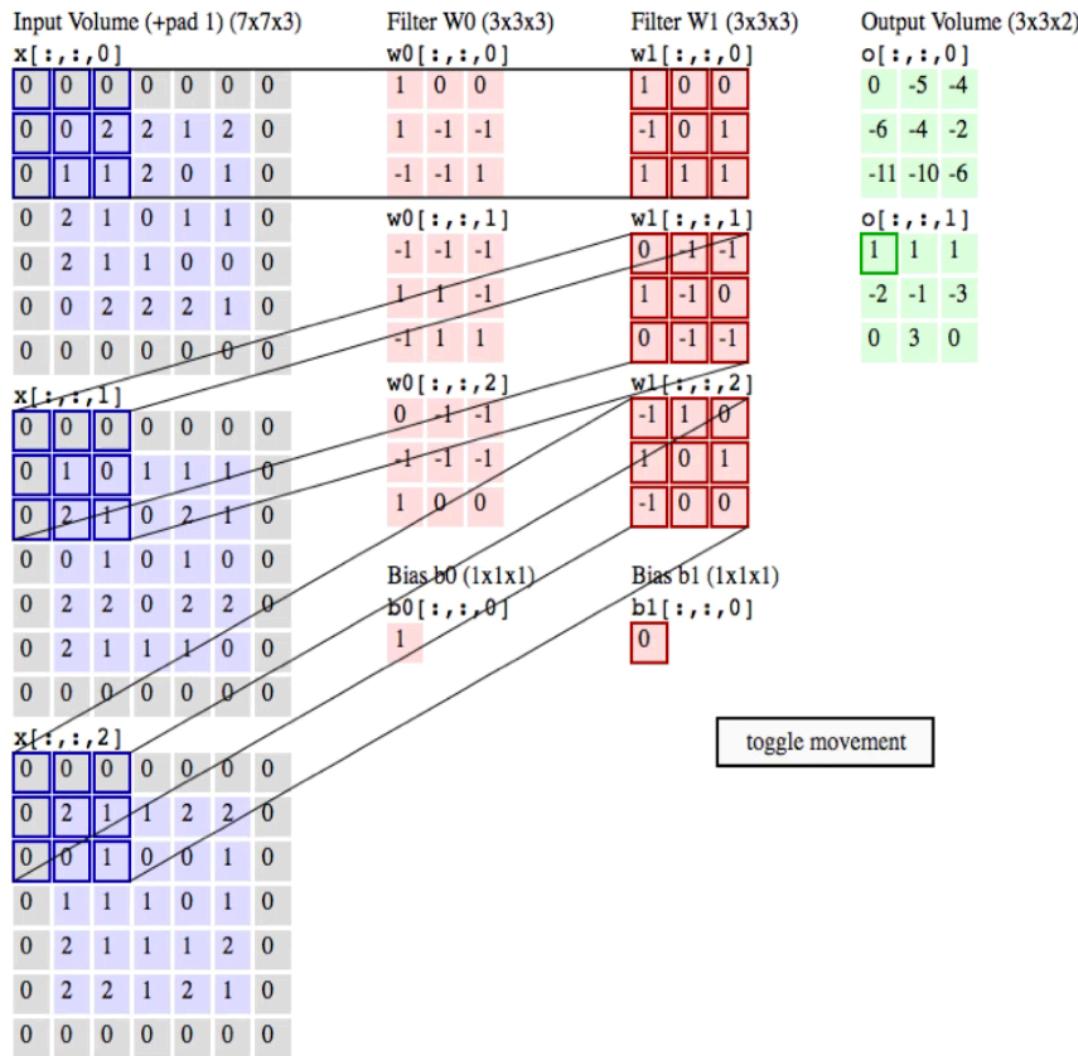
Pooling 1

Convolution 2

Pooling 2

**Fully connected
feed-forward**

What's happening here? (example)



- CNN takes a square of patches of pixels (vs. individual pixels)
- Filters them into a smaller square matrix AND equal to size of the patch (or kernel)
- Job of filter is to find patterns in the pixels

CNN Structure: Convolution 2

Convolution layer 2 maps 32 feature maps to 64 using ReLU (rectifier) activation function

```
# Second convolutional layer  
# maps 32 feature maps to 64.  
W_conv2 = weight_variable([5, 5, 32, 64])  
b_conv2 = bias_variable([64])  
h_conv2 = tf.nn.relu(conv2d(h_pool1,  
W_conv2) + b_conv2)
```

Convolution 1

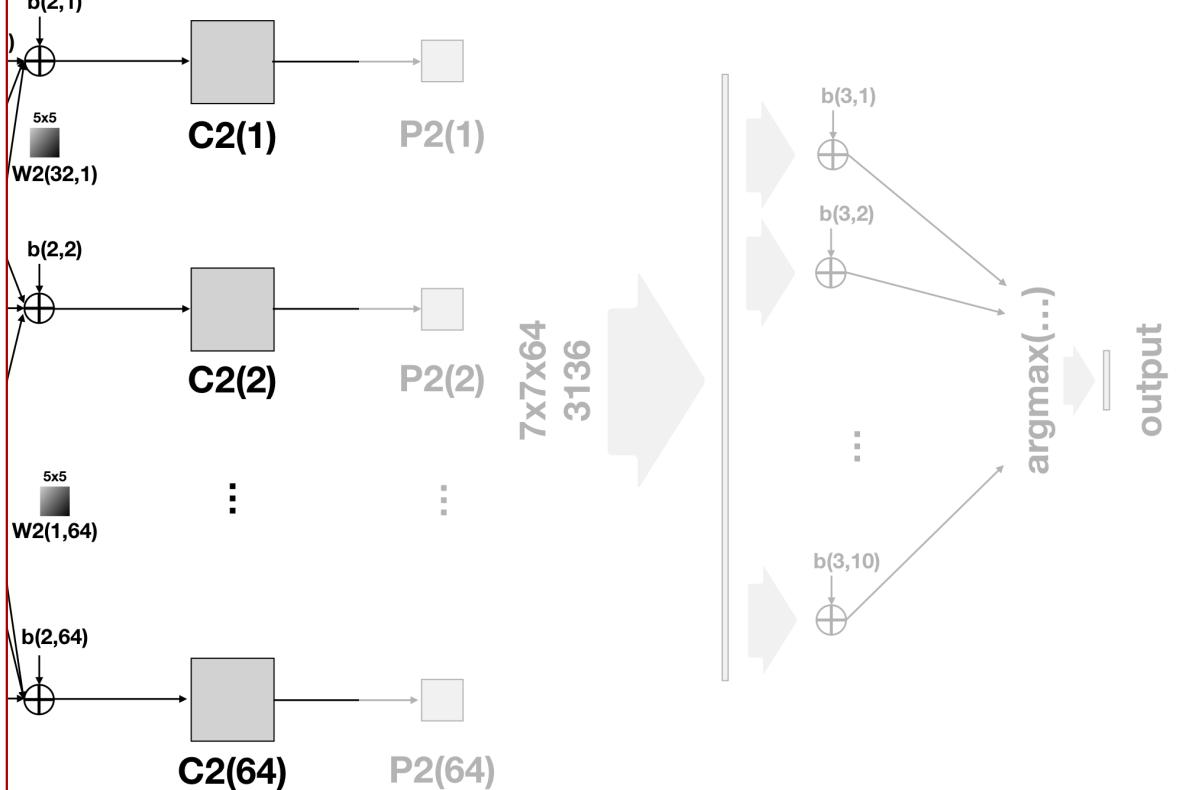
Pooling 1

Convolution 2

Pooling 2

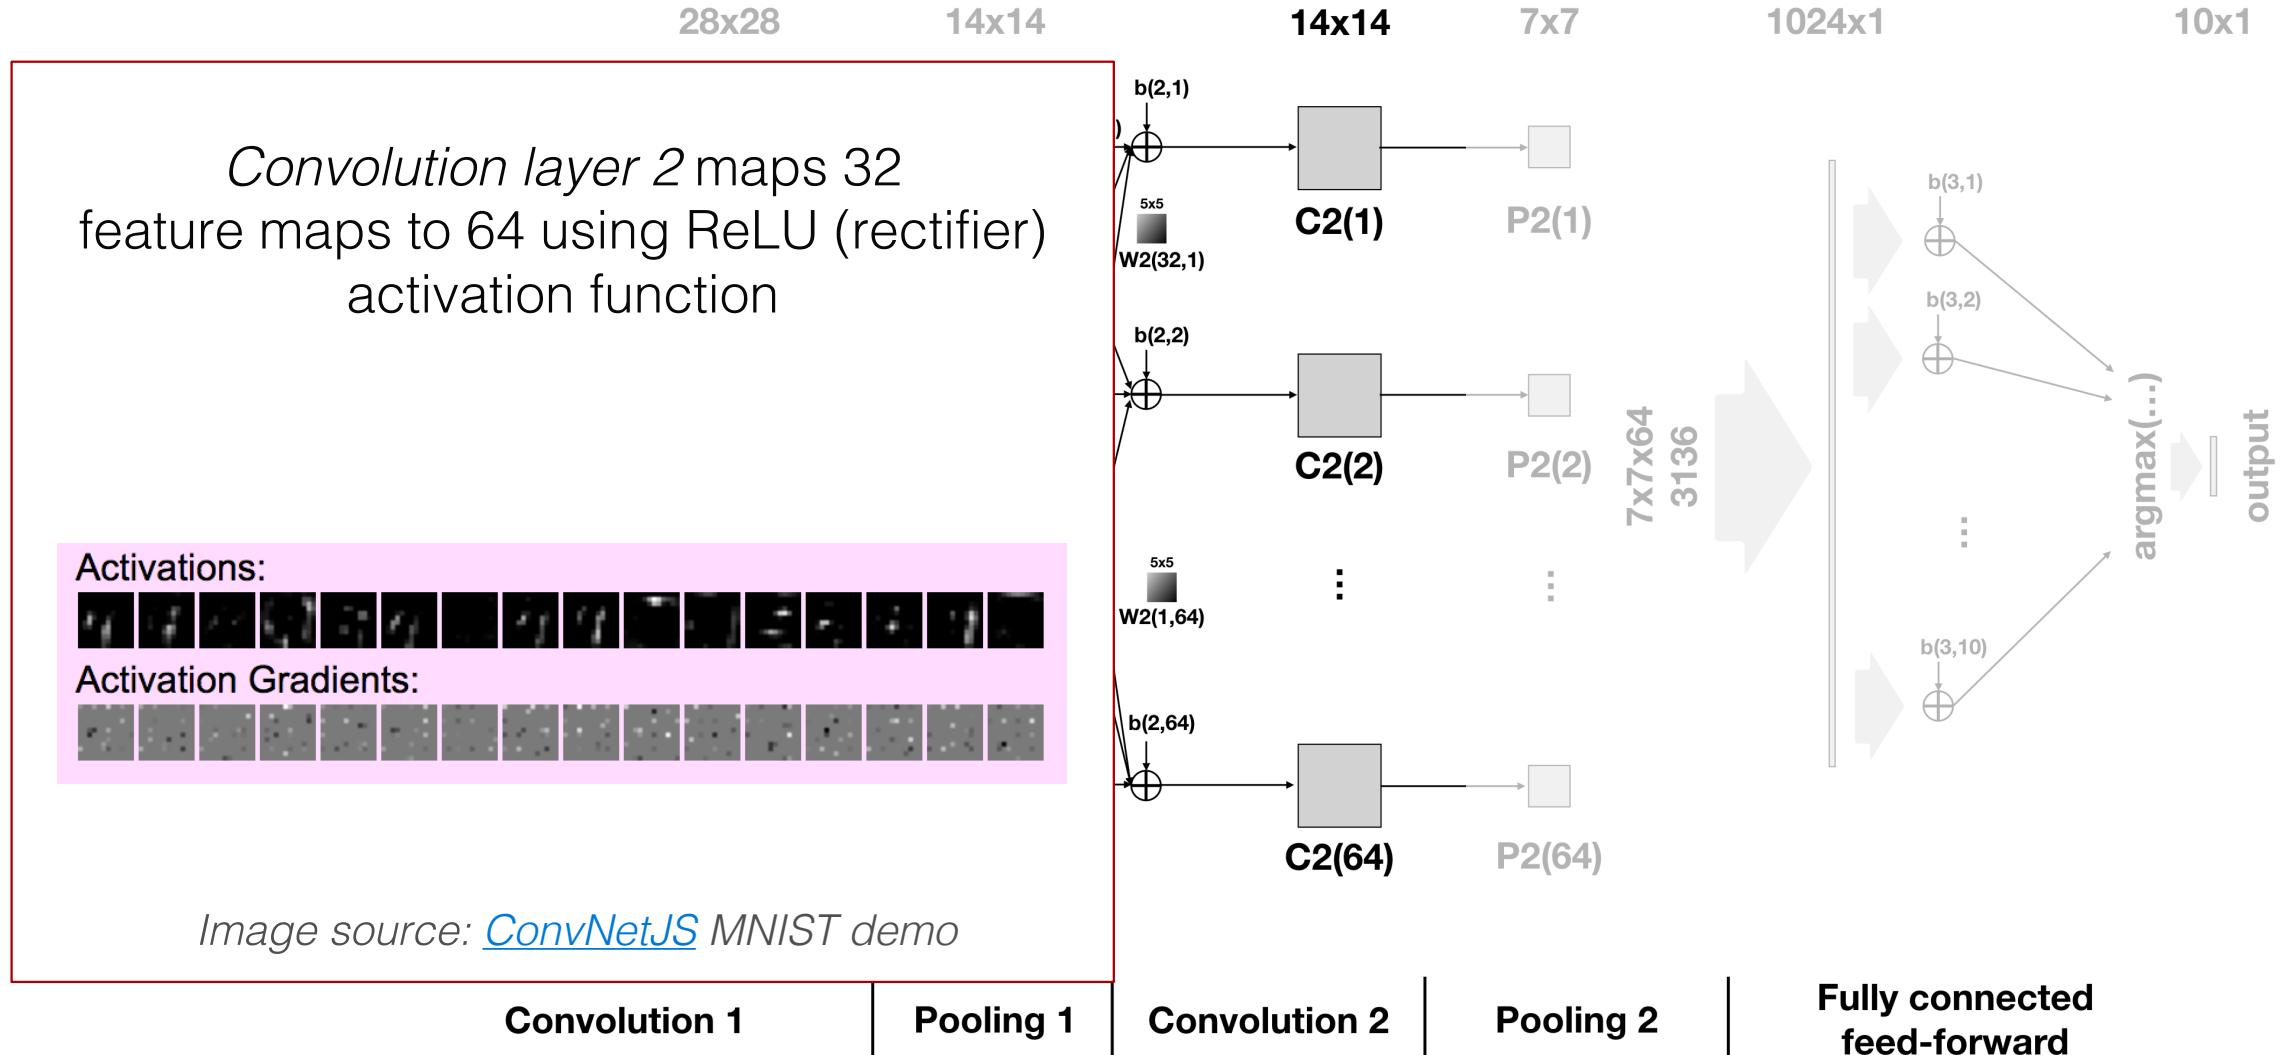
Fully connected feed-forward

28x28 14x14 14x14 7x7 1024x1 10x1

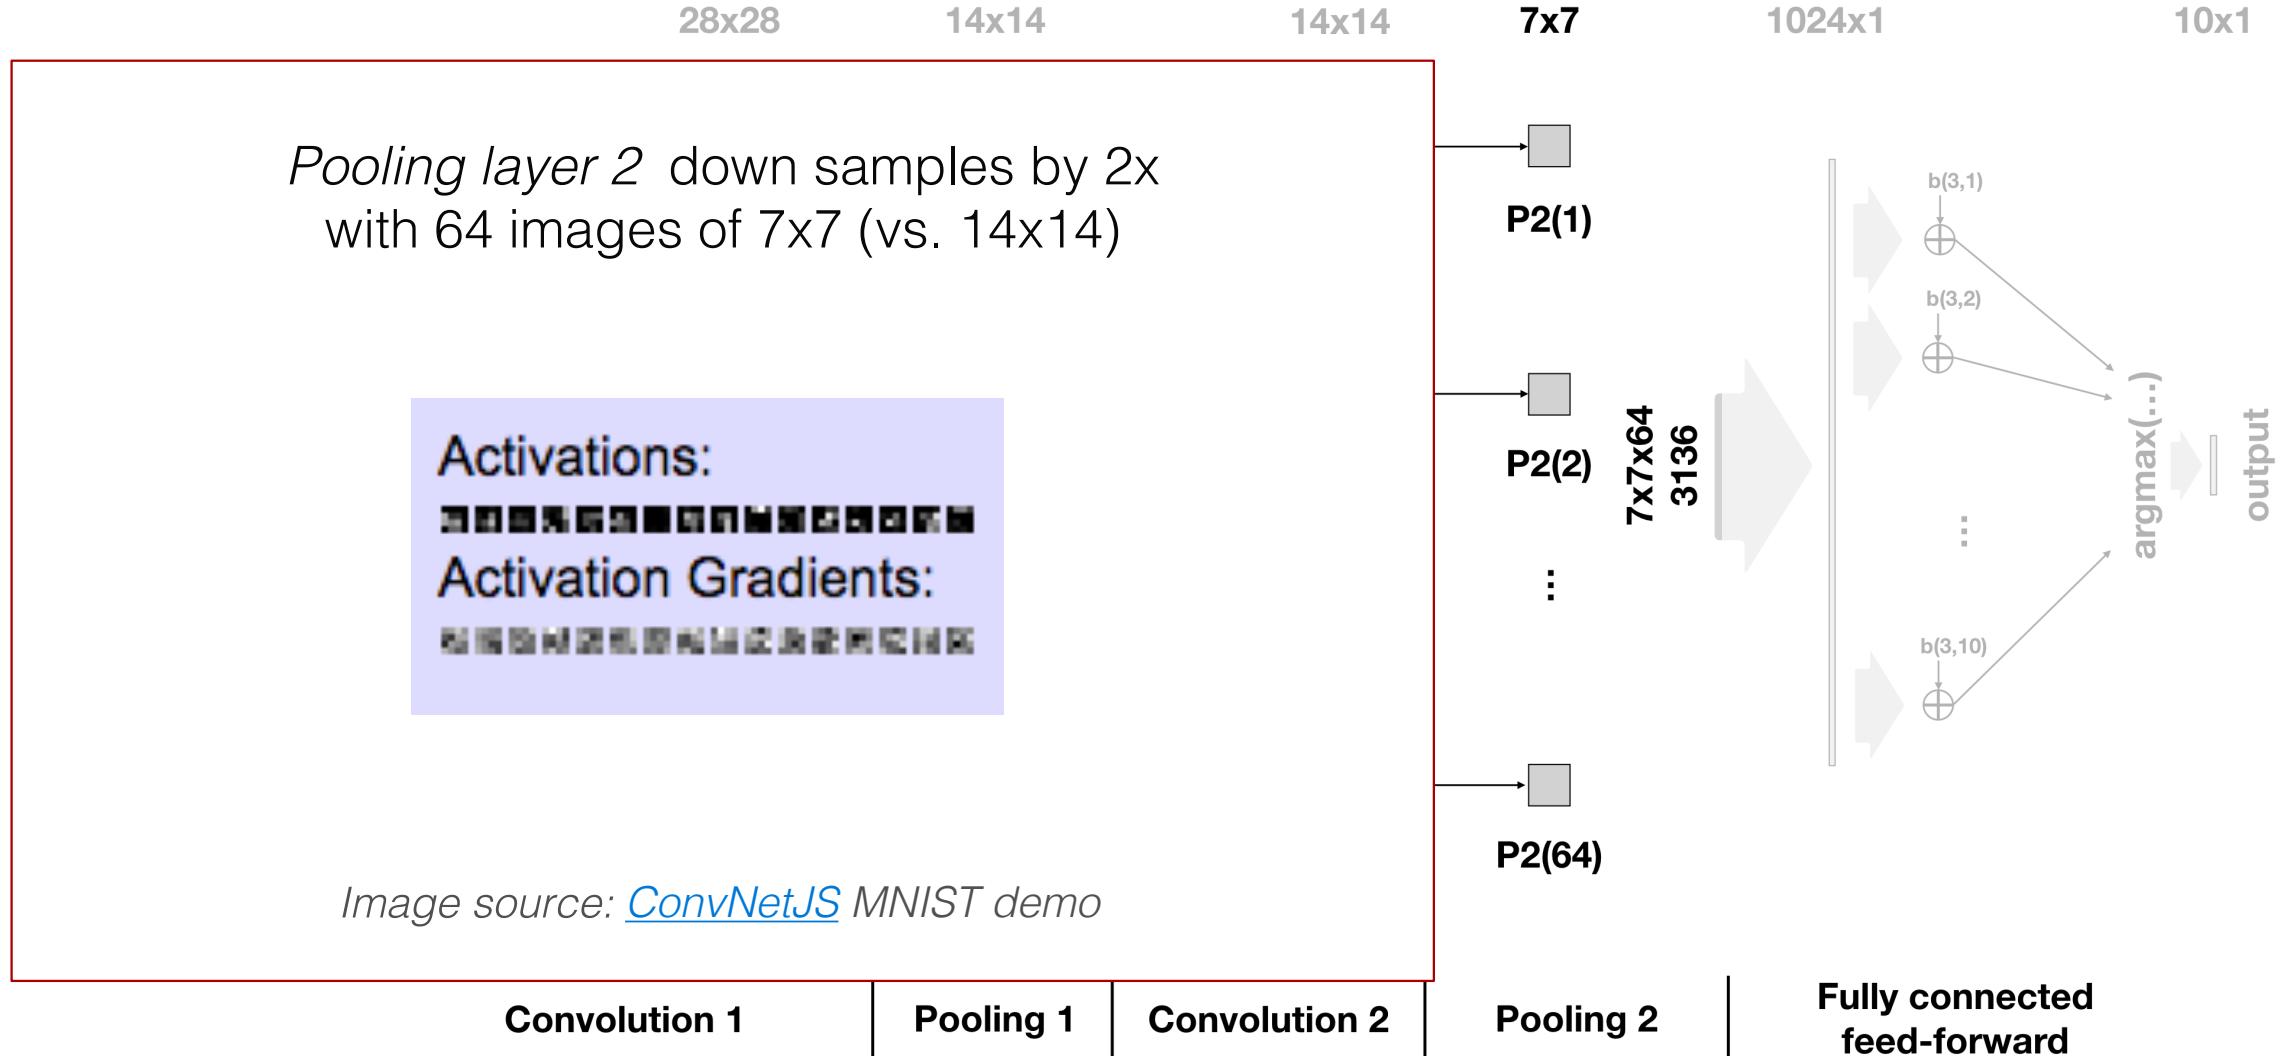


argmax(...)
output

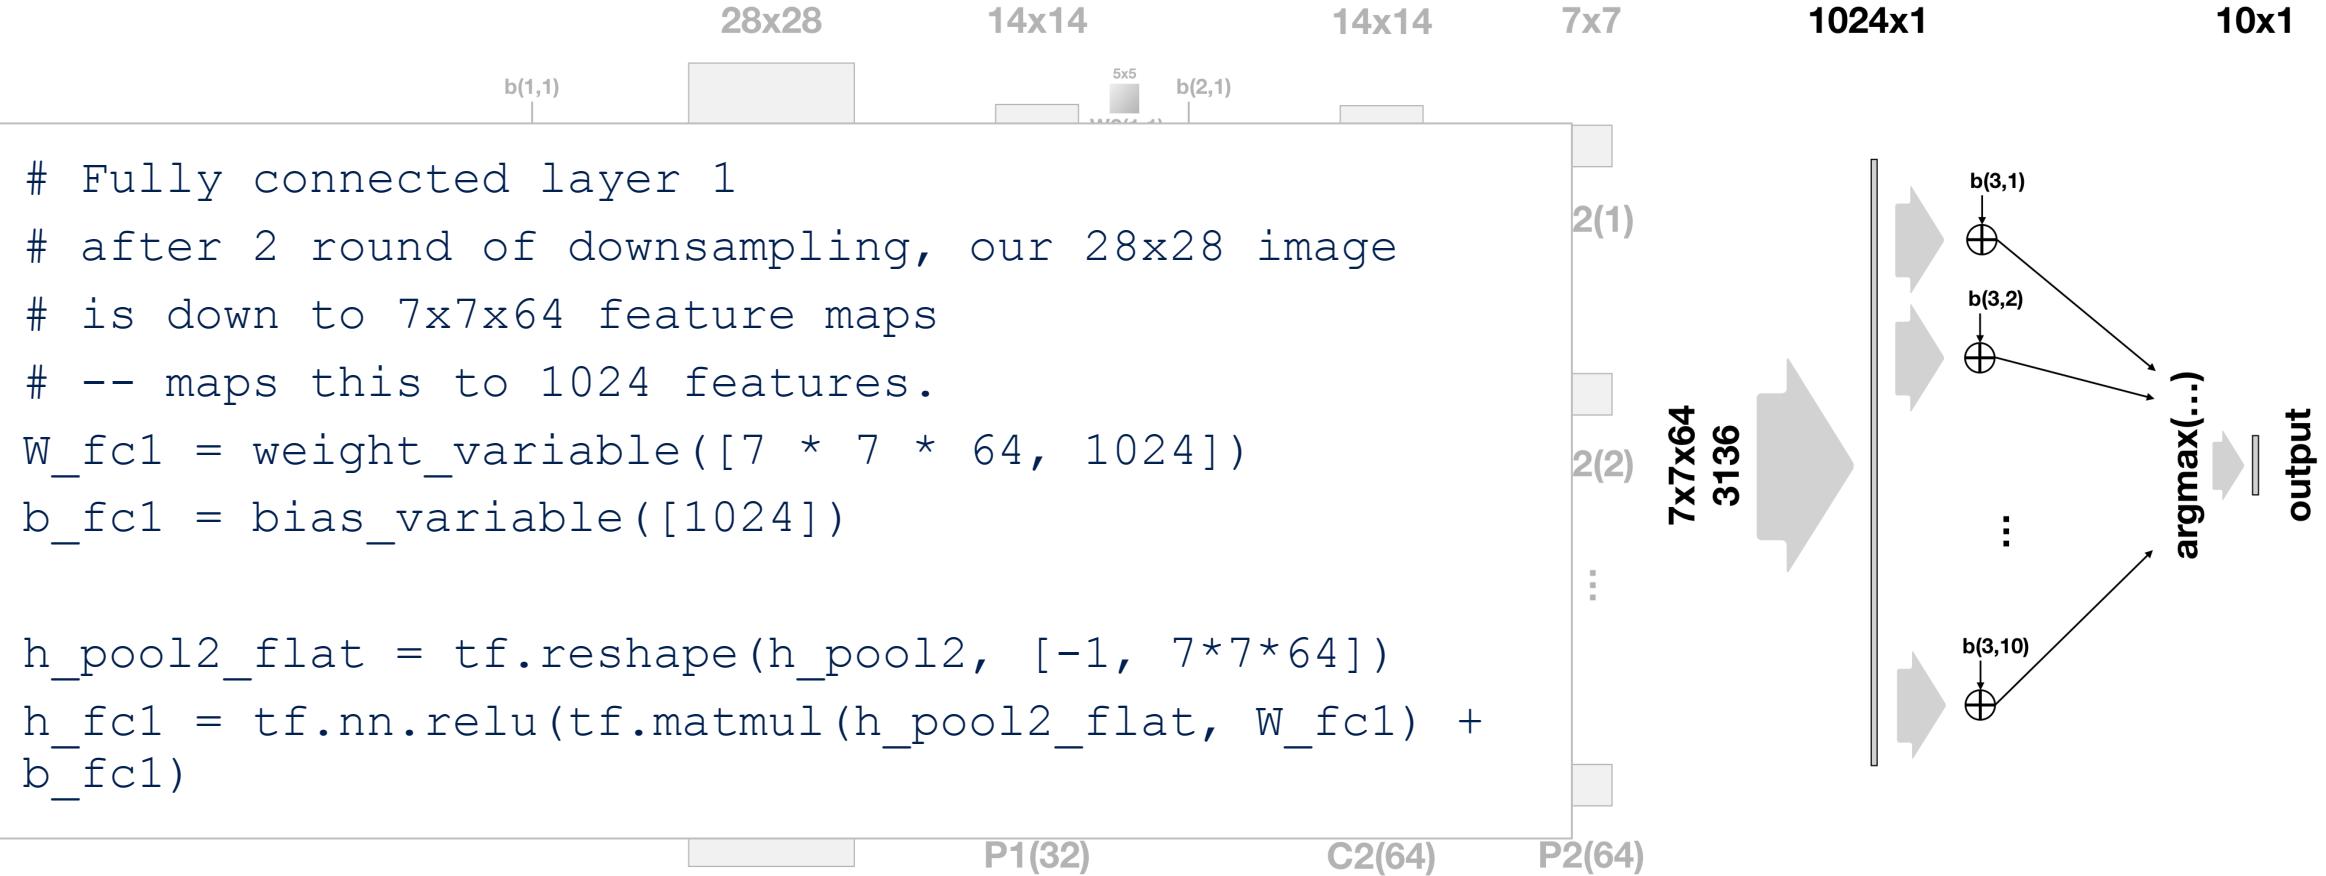
CNN Structure: Convolution 2



CNN Structure: Pooling 2



CNN Structure: Fully connected feed-forward



Convolution 1

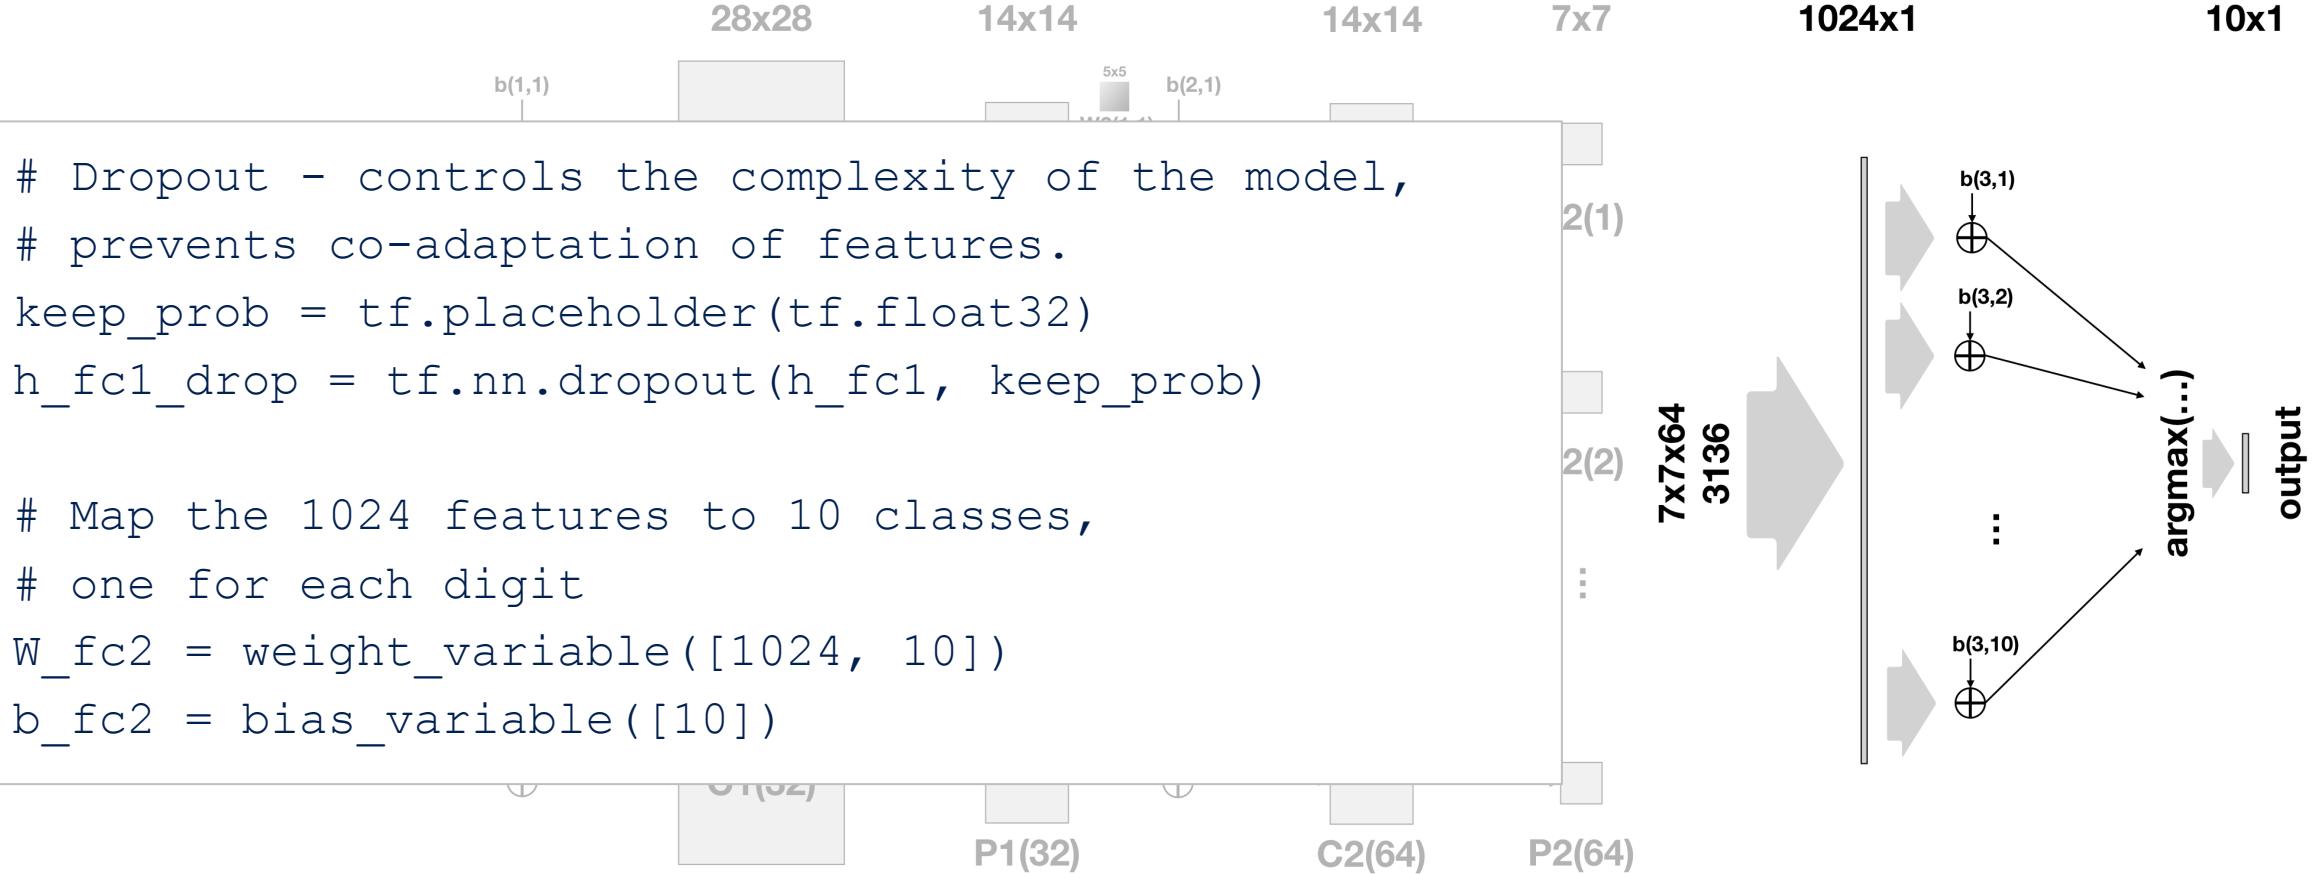
Pooling 1

Convolution 2

Pooling 2

Fully connected
feed-forward

CNN Structure: Dropout



Convolution 1

Pooling 1

Convolution 2

Pooling 2

Fully connected
feed-forward

CNN accuracy: ~98.95

```
1 # Take first 500 images  
2 print('test accuracy %g' % accuracy.eval(feed_dict={x: mnist.test.images[:500, :784], y_: mnist.test.labels[:500,:10], keep_prob: 1.0}))
```

test accuracy 0.966

Command took 0.32 seconds -- by denny.g.lee@gmail.com at 6/26/2017, 8:41:30 PM on unknown cluster

Cmd 8

In fact, the accuracy could be much higher if we set it to higher iterations (e.g. 20000) and tested it against the full test dataset

```
step 19500, training accuracy 1  
step 19600, training accuracy 1  
step 19700, training accuracy 1  
step 19800, training accuracy 1  
step 19900, training accuracy 1  
test accuracy 0.9895
```

Great References

[Andrej Karpathy's ConvNetJS MNIST Demo](#)

[What is back propagation in neural networks?](#)

CS231n: Convolutional Neural Networks for Visual Recognition

[Syllabus and Slides](#) | [Course Notes](#) | [YouTube](#)

- With particular focus on [CS231n: Lecture 7: Convolution Neural Networks](#)

[Neural Networks and Deep Learning](#)

[TensorFlow](#)

Great References

[Deep Visualization Toolbox](#)

[Back Propagation with TensorFlow](#)

[TensorFrames: Google TensorFlow with Apache Spark](#)

[Integrating deep learning libraries with Apache Spark](#)

Upcoming webinar: [Build, Scale, and Deploy Deep Learning Pipelines with Ease](#)



Q&A